

# **Multiagent Reinforcement Learning (MARL)**

September 27, 2013 - ECML'13

---

## **Presenters**

---

- ▶ Daan Bloembergen
- ▶ Daniel Hennes
- ▶ Michael Kaisers
- ▶ Peter Vrancx

# Schedule

---

- ▶ Fundamentals of multi-agent reinforcement learning
  - ▶ **15:30 - 17:00**, Daan Bloembergen and Daniel Hennes
- ▶ Dynamics of learning in strategic interactions
  - ▶ **17:15 - 17:45**, Michael Kaisers
- ▶ Scaling multi-agent reinforcement learning
  - ▶ **17:45 - 18:45**, Peter Vrancx

---

September 27, 2013 - ECML MARL Tutorial

## Who are you?

---

**We would like to get to know our audience!**

---

September 27, 2013 - ECML MARL Tutorial

# Fundamentals of Multi-Agent Reinforcement Learning

Daan Bloembergen and Daniel Hennes

## Outline (1)

---

### Single Agent Reinforcement Learning

- ▶ Markov Decision Processes
  - ▶ Value Iteration
  - ▶ Policy Iteration
- ▶ Algorithms
  - ▶ Q-Learning
  - ▶ Learning Automata

# Outline (2)

---

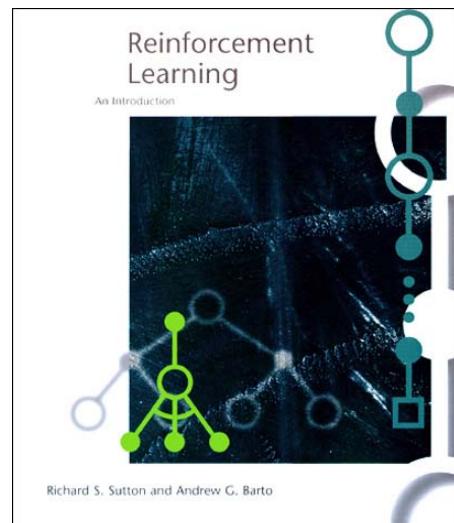
## Multiagent Reinforcement Learning

- ▶ Game Theory
- ▶ Markov Games
  - ▶ Value Iteration
- ▶ Algorithms
  - ▶ Minimax-Q Learning
  - ▶ Nash-Q Learning
  - ▶ Other Equilibrium Learning Algorithms
  - ▶ Policy Hill-Climbing

## Part I: Single Agent Reinforcement Learning

Richard S. Sutton and Andrew G. Barto  
**Reinforcement Learning: An Introduction**  
MIT Press, 1998

Available on-line for free!



## Why reinforcement learning?

Based on ideas from psychology

- ▶ Edward Thorndike's **law of effect**
  - ▶ Satisfaction strengthens behavior, discomfort weakens it
- ▶ B.F. Skinner's **principle of reinforcement**
  - ▶ Skinner Box: train animals by providing (positive) feedback

Learning by interacting with the environment



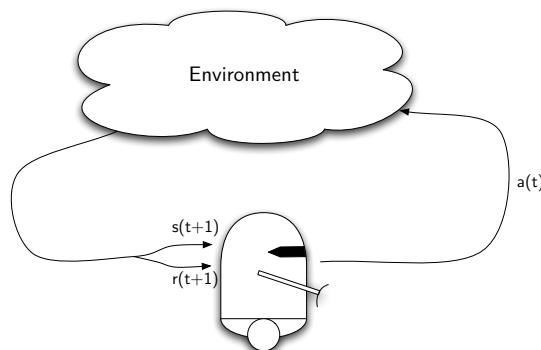
# Why reinforcement learning?

## Control theory

- ▶ Design a controller to minimize some measure of a dynamical systems's behavior
- ▶ Richard Bellman
  - ▶ Use system state and value functions (optimal return)
  - ▶ **Bellman equation**
- ▶ Dynamic programming
  - ▶ Solve optimal control problems by solving the Bellman equation

These two threads came together in the 1980s, producing the modern field of reinforcement learning

# The RL setting



- ▶ Learning from interactions
- ▶ Learning what to do - **how to map situations to actions** - so as to maximize a numerical reward signal

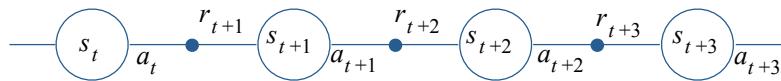
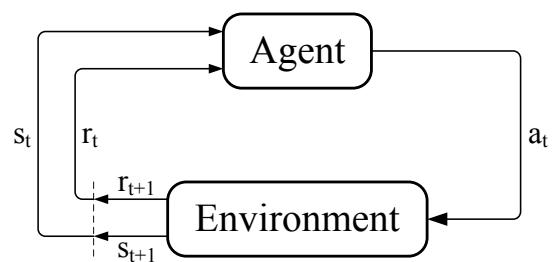
# Key features of RL

- ▶ Learner is **not** told which action to take
- ▶ Trial-and-error approach
- ▶ Possibility of **delayed reward**
  - ▶ Sacrifice short-term gains for greater long-term gains
- ▶ Need to balance **exploration** and **exploitation**
- ▶ In between **supervised** and **unsupervised** learning

## The agent-environment interface

Agent interacts at discrete time steps  $t = 0, 1, 2, \dots$

- ▶ Observes state  $s_t \in S$
- ▶ Selects action  $a_t \in A(s_t)$
- ▶ Obtains immediate reward  $r_{t+1} \in \mathfrak{R}$
- ▶ Observes resulting state  $s_{t+1}$



# Elements of RL

---

- ▶ Time steps need not refer to fixed intervals of real time
- ▶ **Actions** can be
  - ▶ low level (voltage to motors)
  - ▶ high level (go left, go right)
  - ▶ "mental" (shift focus of attention)
- ▶ **States** can be
  - ▶ low level "sensations" (temperature,  $(x, y)$  coordinates)
  - ▶ high level abstractions, symbolic
  - ▶ subjective, internal ("surprised", "lost")
- ▶ The **environment** is not necessarily known to the agent

# Elements of RL

---

- ▶ **State transitions** are
  - ▶ changes to the internal state of the agent
  - ▶ changes in the environment as a result of the agent's action
  - ▶ can be nondeterministic
- ▶ **Rewards** are
  - ▶ goals, subgoals
  - ▶ duration
  - ▶ ...

# Learning how to behave

---

- ▶ The agent's **policy**  $\pi$  at time  $t$  is
  - ▶ a mapping from states to action probabilities
  - ▶  $\pi_t(s, a) = P(a_t = a | s_t = s)$
- ▶ Reinforcement learning methods specify **how** the agent changes its policy as a result of experience
- ▶ Roughly, the agent's goal is to get **as much reward** as it can over the long run

## The objective

---

Suppose the sequence of rewards after time  $t$  is

$$r_{t+1}, r_{t+2}, r_{t+3}, \dots$$

- ▶ The goal is to maximize the **expected return**  $E\{R_t\}$  for each time step  $t$
- ▶ **Episodic tasks** naturally break into episodes, e.g., plays of a game, trips through a maze

$$R_t = r_{t+1} + r_{t+2} + \dots + r_T$$

# The objective

---

- ▶ **Continuing tasks** do not naturally break up into episodes
- ▶ Use **discounted return** instead of total reward

$$R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

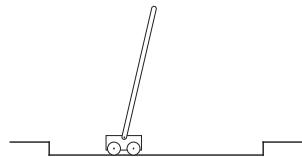
where  $\gamma$ ,  $0 \leq \gamma \leq 1$  is the **discount factor** such that

shortsighted  $0 \leftarrow \gamma \rightarrow 1$  farsighted

## Example: pole balancing

---

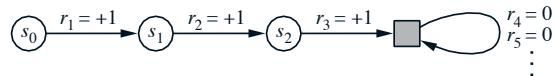
- ▶ As an **episodic** task where each episode ends upon failure
  - ▶ reward = +1 for each step before failure
  - ▶ return = number of steps before failure
- ▶ As a **continuing** task with discounted return
  - ▶ reward = -1 upon failure
  - ▶ return =  $-\gamma^k$ , for  $k$  steps before failure
- ▶ In both cases, return is maximized by avoiding failure for as long as possible



# A unified notation

---

- ▶ Think of each episode as ending in an **absorbing state** that always produces a reward of zero



- ▶ Now we can cover both episodic and continuing tasks by writing

$$R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

## Markov decision processes

---

- ▶ It is often useful to assume that all relevant information is present in the current state: **Markov property**

$$P(s_{t+1}, r_{t+1} | s_t, a_t) = P(s_{t+1}, r_{t+1} | s_t, a_t, r_t, s_{t-1}, a_{t-1}, \dots, r_1, s_0, a_0)$$

- ▶ If a reinforcement learning task has the Markov property, it is basically a **Markov Decision Process (MDP)**
- ▶ Assuming finite state and action spaces, it is a finite MDP

# Markov decision processes

---

An MDP is defined by

- ▶ **State and action sets**
- ▶ One-step dynamics defined by state **transition probabilities**

$$\mathcal{P}_{ss'}^a = P(s_{t+1} = s' | s_t = s, a_t = a)$$

- ▶ **Reward probabilities**

$$\mathcal{R}_{ss'}^a = E(r_{t+1} | s_t = s, a_t = a, s_{t+1} = s')$$

## Value functions

---

- ▶ When following a fixed policy  $\pi$  we can define the **value** of a state  $s$  under that policy as

$$V^\pi(s) = E_\pi(R_t | s_t = s) = E_\pi\left(\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s\right)$$

- ▶ Similarly we can define the value of taking action  $a$  in state  $s$  as

$$Q^\pi(s, a) = E_\pi(R_t | s_t = s, a_t = a)$$

# Value functions

---

- ▶ The value function has a particular recursive relationship, defined by the **Bellman equation**

$$V^\pi(s) = \sum_a \pi(s, a) \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V^\pi(s')]$$

- ▶ The equation expresses the recursive relation between the value of a state and its successor states, and averages over all possibilities, weighting each by its probability of occurring

# Optimal policy for an MDP

---

- ▶ We want to find the policy that maximizes long term reward, which equates to finding the optimal value function

$$V^*(s) = \max_\pi V^\pi(s) \quad \forall s \in S$$

$$Q^*(s, a) = \max_\pi Q^\pi(s, a) \quad \forall s \in S, a \in A(s)$$

- ▶ Expressed recursively, this is the **Bellman optimality equation**

$$\begin{aligned} V^*(s) &= \max_{a \in A(s)} Q^{*\pi}(s, a) \\ &= \max_{a \in A(s)} \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V^*(s')] \end{aligned}$$

# Solving the Bellman equation

- ▶ We can find the **optimal policy** by solving the Bellman equation
  - ▶ Dynamic Programming
- ▶ Two approaches:
  - ▶ Iteratively improve the value function: **value iteration**
  - ▶ Iteratively evaluate and improve the policy: **policy iteration**
- ▶ Both approaches are proven to converge to the optimal value function

## Value iteration

```
Initialize  $V$  arbitrarily, e.g.,  $V(s) = 0$ , for all  $s \in \mathcal{S}^+$ 
Repeat
     $\Delta \leftarrow 0$ 
    For each  $s \in \mathcal{S}$ :
         $v \leftarrow V(s)$ 
         $V(s) \leftarrow \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$ 
         $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
    until  $\Delta < \theta$  (a small positive number)

    Output a deterministic policy,  $\pi$ , such that
    
$$\pi(s) = \arg \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$$

```

# Policy iteration

- ▶ Often the optimal policy has been reached long before the value function has converged
- ▶ Policy iteration calculates a new policy **based on the current value function**, and then calculates a new value function based on this policy
- ▶ This process often converges faster to the optimal policy

# Policy iteration

```
1. Initialization  
     $V(s) \in \Re$  and  $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$   
  
2. Policy Evaluation  
    Repeat  
         $\Delta \leftarrow 0$   
        For each  $s \in \mathcal{S}$ :  
             $v \leftarrow V(s)$   
             $V(s) \leftarrow \sum_{s'} \mathcal{P}_{ss'}^{\pi(s)} [\mathcal{R}_{ss'}^{\pi(s)} + \gamma V(s')]$   
             $\Delta \leftarrow \max(\Delta, |v - V(s)|)$   
    until  $\Delta < \theta$  (a small positive number)  
  
3. Policy Improvement  
     $policy-stable \leftarrow true$   
    For each  $s \in \mathcal{S}$ :  
         $b \leftarrow \pi(s)$   
         $\pi(s) \leftarrow \arg \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$   
        If  $b \neq \pi(s)$ , then  $policy-stable \leftarrow false$   
    If  $policy-stable$ , then stop; else go to 2
```

# Learning an optimal policy online

---

- ▶ Both previous approaches require to know the dynamics of the environment
- ▶ Often this information is not available
- ▶ Using **temporal difference (TD)** methods is one way of overcoming this problem
  - ▶ Learn directly from raw experience
  - ▶ No model of the environment required (model-free)
  - ▶ E.g.: **Q-learning**
- ▶ Update predicted state values based on new observations of immediate rewards and successor states

## Q-learning

---

- ▶ Q-learning updates state-action values based on the immediate reward and the optimal expected return

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[ r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

- ▶ Directly learns the optimal value function independent of the policy being followed
  - ▶ In contrast to on-policy learners, e.g. **SARSA**
- ▶ Proven to converge to the optimal policy given "sufficient" updates for each state-action pair, and decreasing learning rate  $\alpha$  [Watkins92]

# Q-learning

```
Initialize  $Q(s, a)$  arbitrarily
Repeat (for each episode):
    Initialize  $s$ 
    Repeat (for each step of episode):
        Choose  $a$  from  $s$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)
        Take action  $a$ , observe  $r, s'$ 
         $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ 
         $s \leftarrow s'$ ;
    until  $s$  is terminal
```

## Action selection

- ▶ How to select an action based on the values of the states or state-action pairs?
- ▶ Success of RL depends on a **trade-off**
  - ▶ Exploration
  - ▶ Exploitation
- ▶ **Exploration** is needed to prevent getting stuck in local optima
- ▶ To ensure convergence you need to **exploit**

# Action selection

---

Two common choices

- ▶  **$\epsilon$ -greedy**
  - ▶ Choose the best action with probability  $1 - \epsilon$
  - ▶ Choose a random action with probability  $\epsilon$
- ▶ **Boltzmann exploration** (softmax) uses a temperature parameter  $\tau$  to balance exploration and exploitation

$$\pi_t(s, a) = \frac{e^{Q_t(s, a)/\tau}}{\sum_{a' \in A} e^{Q_t(s, a')/\tau}}$$

pure exploitation  $0 \leftarrow \tau \rightarrow \infty$  pure exploration

# Learning automata

---

- ▶ **Learning automata** [Narendra74] directly modify their policy based on the observed reward (policy iteration)
- ▶ Finite action-set learning automata learn a policy over a finite set of actions

$$\pi'(a) = \pi(a) + \begin{cases} \alpha r(1 - \pi(a)) - \beta(1 - r)\pi(a) & \text{if } a = a_t \\ -\alpha r\pi(a) + \beta(1 - r)[(k-1)^{-1} - \pi(a)] & \text{if } a \neq a_t \end{cases}$$

where  $k = |A|$ , and  $\alpha$  and  $\beta$  are reward and penalty parameters respectively, and  $r \in [0, 1]$

- ▶ **Cross learning** is a special case where  $\alpha = 1$  and  $\beta = 0$

# Networks of learning automata

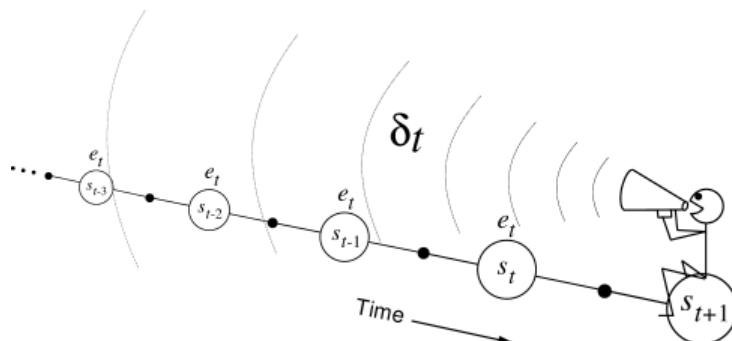
- ▶ A single learning automaton ignores any state information
- ▶ In a **network of learning automata** [Wheeler86] control is passed on from one automaton to another
  - ▶ One automaton  $\mathcal{A}$  is active for each state
  - ▶ The immediate reward  $r$  is replaced by the average cumulative reward  $\bar{r}$  since the last visit to that state

$$\bar{r}_t(s) = \frac{\Delta r}{\Delta t} = \frac{\sum_{i=l(s)}^{t-1} r_i}{t - l(s)}$$

where  $l(s)$  indicates in which time step state  $s$  was last visited

## Extensions

- ▶ Multi-step TD: **eligibility traces**
  - ▶ Instead of observing one immediate reward, use  $n$  consecutive rewards for the value update
  - ▶ Intuition: your current choice of action may have implications for the future
  - ▶ State-action pairs are eligible for future rewards, with more recent states getting more credit



# Extensions

---

## ► Reward shaping

- ▶ Incorporate domain knowledge to provide additional rewards during an episode
- ▶ Guide the agent to learn faster
- ▶ (Optimal) policies preserved given a potential-based shaping function [Ng99]

## ► Function approximation

- ▶ So far we have used a tabular notation for value functions
- ▶ For large state and actions spaces this approach becomes intractable
- ▶ Function approximators can be used to generalize over large or even continuous state and action spaces

# Questions so far?

---



# Part II: Multiagent Reinforcement Learning

## Preliminaries: Fundamentals of Game Theory

## Game theory

- ▶ Models **strategic interactions** as games
- ▶ In **normal form games**, all players simultaneously select an action, and their joint action determines their individual payoff
  - ▶ One-shot interaction
  - ▶ Can be represented as an  $n$ -dimensional payoff matrix, for  $n$  players
- ▶ A player's **strategy** is defined as a probability distribution over his possible actions

# Example: Prisoner's Dilemma

---

- ▶ Two prisoners (A and B) commit a crime together
- ▶ They are questioned separately and can choose to confess or deny
  - ▶ If both confess, both prisoners will serve 3 years in jail
  - ▶ If both deny, both serve only 1 year for minor charges
  - ▶ If only one confesses, he goes free, while the other serves 5 years

	C	D
C	-3, -3	-0, -5
D	-5, -0	-1, -1

# Example: Prisoner's Dilemma

---

- ▶ What should they do?
- ▶ If both deny, their total penalty is lowest
  - ▶ But is this individually rational?
- ▶ Purely selfish: regardless of what the other player does, confess is the optimal choice
  - ▶ If the other confesses, 3 instead of 5 years
  - ▶ If the other denies, free instead of 1 year

	C	D
C	-3, -3	-0, -5
D	-5, -0	-1, -1

# Solution concepts

## ► Nash equilibrium

- ▶ Individually rational
- ▶ No player can improve by unilaterally changing his strategy
- ▶ Mutual confession is the only Nash equilibrium of this game
- ▶ Jointly the players could do better
  - ▶ **Pareto optimum:** there is no other solution for which all players do at least as well and at least one player is strictly better off
  - ▶ Mutual denial Pareto dominates the Nash equilibrium in this game

	C	D
C	-3, -3	-0, -5
D	-5, -0	-1, -1

# Types of games

## ► Competitive or zero-sum

- ▶ Players have opposing preferences
- ▶ E.g. Matching Pennies

**Matching Pennies**

	H	T
H	+1, -1	-1, +1
T	-1, +1	+1, -1

## ► Symmetric games

- ▶ Players are identical
- ▶ E.g. Prisoner's Dilemma

**Prisoner's Dilemma**

	C	D
C	-3, -3	-0, -5
D	-5, -0	-1, -1

## ► Asymmetric games

- ▶ Players are unique
- ▶ E.g. Battle of the Sexes

**Battle of the Sexes**

	B	S
B	2, 1	0, 0
S	0, 0	1, 2

# Part II: Multiagent Reinforcement Learning

## MARL: Motivation

- ▶ MAS offer a solution paradigm that can cope with complex problems
- ▶ Technological challenges require decentralised solutions
  - ▶ Multiple autonomous vehicles for exploration, surveillance or rescue missions
  - ▶ Distributed sensing
  - ▶ Traffic control (data, urban or air traffic)
- ▶ Key advantages: Fault tolerance and load balancing
- ▶ **But:** highly dynamic and nondeterministic environments!
- ▶ Need for adaptation on an individual level
- ▶ **Learning is crucial!**

# MARL: From single to multiagent learning

---

- ▶ Inherently more challenging
- ▶ Agents interact with the environment and each other
- ▶ Learning is simultaneous
- ▶ Changes in strategy of one agent might affect strategy of other agents
- ▶ Questions:
  - ▶ One vs. many learning agents?
  - ▶ Convergence?
  - ▶ Objective: maximise common reward or individual reward?
  - ▶ Credit assignment?

## Independent reinforcement learners

---

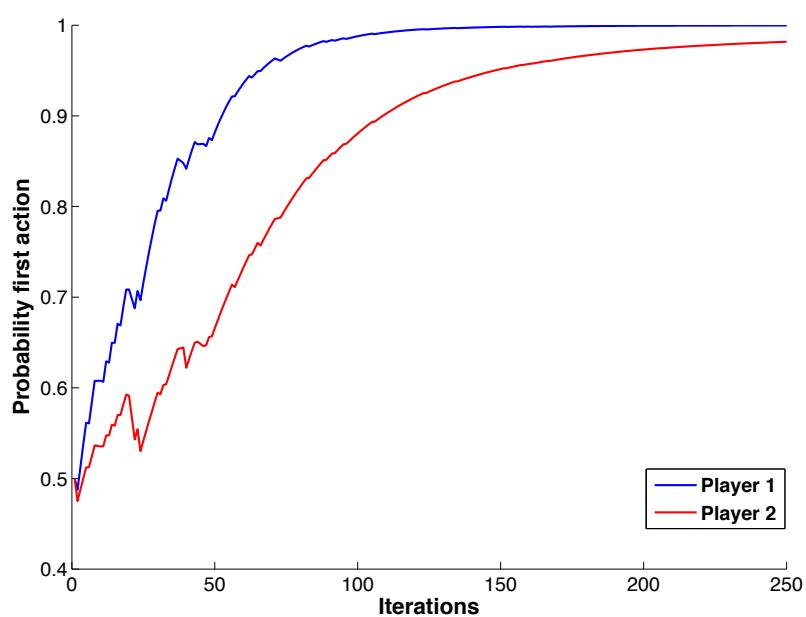
- ▶ Naive extension to multi agent setting
- ▶ Independent learners mutually ignore each other
- ▶ Implicitly perceive interaction with other agents as noise in a stochastic environment

# Learning in matrix games

- ▶ Two Q-learners interact in Battle of the Sexes
  - ▶  $\alpha = 0.01$
  - ▶ Boltzmann exploration with  $\tau = 0.2$
- ▶ They only observe their immediate reward
- ▶ Policy is gradually improved

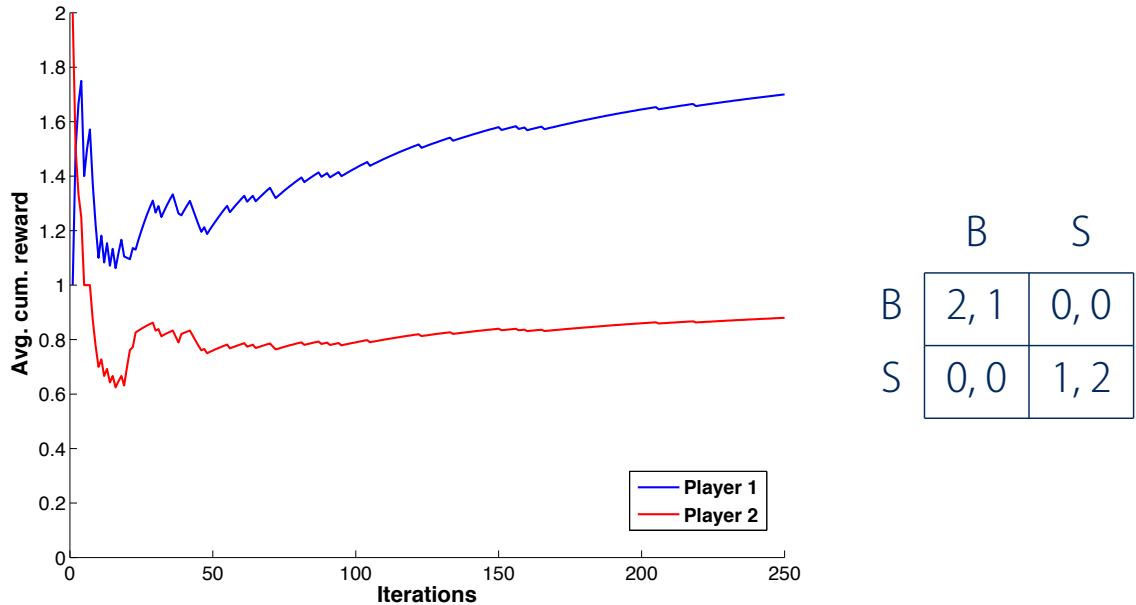
	B	S
B	2, 1	0, 0
S	0, 0	1, 2

# Learning in matrix games



	B	S
B	2, 1	0, 0
S	0, 0	1, 2

# Learning in matrix games



## Markov games

$n$ -player game:  $\langle n, S, A^1, \dots, A^n, \mathcal{R}^1, \dots, \mathcal{R}^n, \mathcal{P} \rangle$

- ▶  $S$ : set of states
- ▶  $A^i$ : action set for player  $i$
- ▶  $\mathcal{R}^i$ : reward/payoff for player  $i$
- ▶  $\mathcal{P}$ : transition function

The payoff function  $\mathcal{R}^i : S \times A^1 \times \dots \times A^n \mapsto \mathbb{R}$  maps the joint action  $a = \langle a^1 \dots a^n \rangle$  to an immediate payoff value for player  $i$ .

The transition function  $\mathcal{P} : S \times A^1 \times \dots \times A^n \mapsto \Delta(S)$  determines the probabilistic state change to the next state  $s_{t+1}$ .

# Value iteration in Markov games

---

Single agent MDP:

$$\begin{aligned} V^*(s) &= \max_{a \in A(s)} Q^{\pi^*}(s, a) \\ &= \max_{a \in A(s)} \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V^*(s')] \end{aligned}$$

2-player zero-sum stochastic game:

$$Q^*(s, \langle a^1, a^2 \rangle) = \mathcal{R}(s, \langle a^1, a^2 \rangle) + \gamma \sum_{s' \in S} \mathcal{P}_{s'}(s, \langle a^1, a^2 \rangle) V^*(s')$$

$$V^*(s) = \max_{\pi \in \Delta(A^1)} \min_{a^2 \in A^2} \sum_{a^1 \in A^1} \pi_{a^1} Q^*(s, \langle a^1, a^2 \rangle)$$

## Minimax-Q

---

- ▶ Value iteration requires knowledge of the reward and transition functions
- ▶ Minimax- $Q$  [Littman94]: learning algorithm for zero-sum games
- ▶ Payoffs balance out, each agent only needs to observe its own payoff
- ▶  $Q$  is a function of the joint action:

$$Q(s, \langle a^1, a^2 \rangle) = \mathcal{R}(s, \langle a^1, a^2 \rangle) + \gamma \sum_{s' \in S} \mathcal{P}_{s'}(s, \langle a^1, a^2 \rangle) V(s')$$

- ▶ A joint action learner (JAL) is an agent that learns  $Q$ -values for joint actions as opposed to individual actions.

## Minimax-Q (2)

Update rule for agent 1 with reward function  $\mathcal{R}_t$  at stage  $t$ :

$$Q_{t+1}(s_t, \langle a_t^1, a_t^2 \rangle) = (1 - \alpha_t) Q_t(s_t, \langle a_t^1, a_t^2 \rangle) + \alpha_t [\mathcal{R}_t + \gamma V_t(s_{t+1})]$$

The value of the next state  $V(s_{t+1})$ :

$$V_{t+1}(s) = \max_{\pi \in \Delta(A^1)} \min_{a^2 \in A^2} \sum_{a^1 \in A^1} \pi_{a^1} Q_t(s, \langle a^1, a^2 \rangle).$$

Minimax- $Q$  converges to Nash equilibria under the same assumptions as regular  $Q$ -learning [Littman94]

## Nash- $Q$ learning

- ▶ Nash- $Q$  learning [Hu03]: joint action learner for general-sum stochastic games
- ▶ Each individual agent has to estimate  $Q$  values for all other agents as well
- ▶ **Optimal Nash- $Q$  values:** sum of immediate reward and discounted future rewards under the condition that all agents play a specified Nash equilibrium from the next stage onward

# Nash-Q learning (2)

---

Update rule for agent  $i$ :

$$Q_{t+1}^i(s_t, \langle a^1, \dots, a^n \rangle) = (1 - \alpha_t) Q(s_t, \langle a^1, \dots, a^n \rangle) + \alpha_t [\mathcal{R}_t + \gamma \text{Nash } V_t^i(s_{t+1})]$$

A Nash equilibrium is computed for each stage game  $(Q_t^1(s_{t+1}, \cdot), \dots, Q_t^n(s_{t+1}, \cdot))$  and results in the equilibrium payoff  $\text{Nash } V_t^i(s_{t+1}, \cdot)$  to agent  $i$

Agent  $i$  uses the same update rule to estimate  $Q$  values for all other agents, i.e.,  $Q^j \forall j \in \{1, \dots, n\} \setminus i$

## Other equilibrium learning algorithms

---

- ▶ Friend-or-Foe  $Q$ -learning [Littman01]
- ▶ Correlated- $Q$  learning (CE- $Q$ ) [Greenwald03]
- ▶ Nash bargaining solution  $Q$ -learning (NBS- $Q$ ) [Qiao06]
- ▶ Optimal adaptive learning (OAL) [Wang02]
- ▶ Asymmetric- $Q$  learning [Kononenko03]

# Limitations of MARL

---

- ▶ Convergence guarantees are mostly restricted to stateless repeated games
- ▶ ... or are inapplicable in general-sum games
- ▶ Many convergence proofs have strong assumptions with respect to a-priori knowledge and/or observability
- ▶ Equilibrium learners focus on stage-wise solutions (only indirect state coupling)

## Summary

---

In a multi-agent system

- ▶ be aware what information is available to the agent
- ▶ if you can afford to try, just run an algorithm that matches the assumptions
- ▶ proofs of convergence are available for small games
- ▶ new research can focus either on engineering solutions, or advancing the state-of-the-art theories

# Questions so far?

---



## Thank you!

Daan Bloembergen | daan.bloembergen@gmail.com  
Daniel Hennes | daniel.hennes@gmail.com

# Dynamics of Learning in Strategic Interactions

Michael Kaisers

## Outline

---

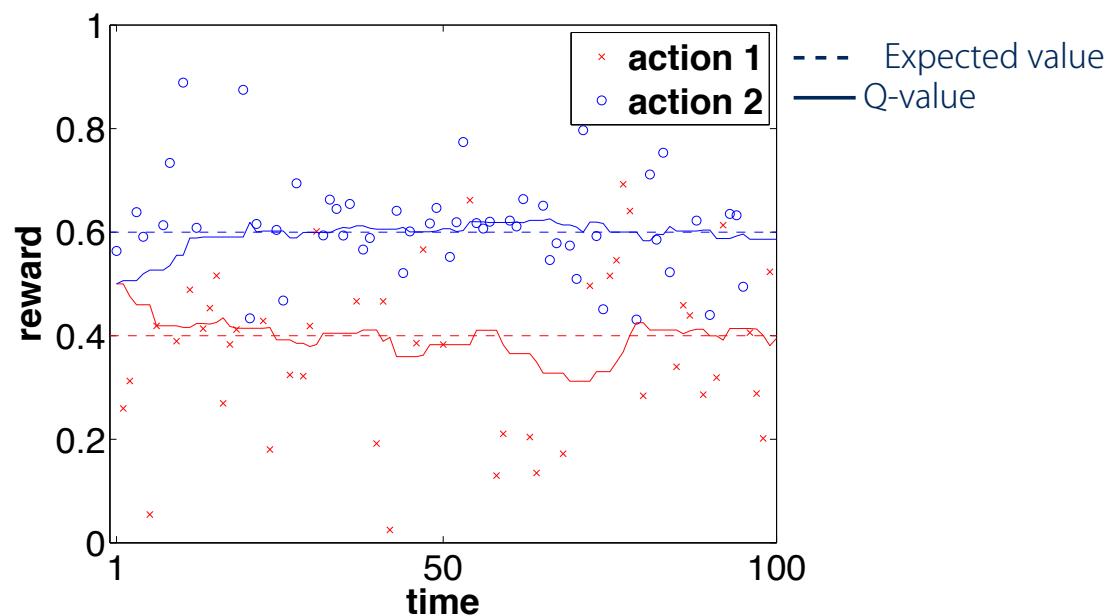
- ▶ Action values in dynamic environments
- ▶ Deriving learning dynamics
- ▶ Illustrating Convergence
- ▶ Comparing dynamics of various algorithms
- ▶ Replicator dynamics as models of evolution, swarm intelligence and learning
- ▶ Summary

# Action values in dynamic environments

Action values are estimated by sampling from interactions with the environment, possibly in the presence of other agents.

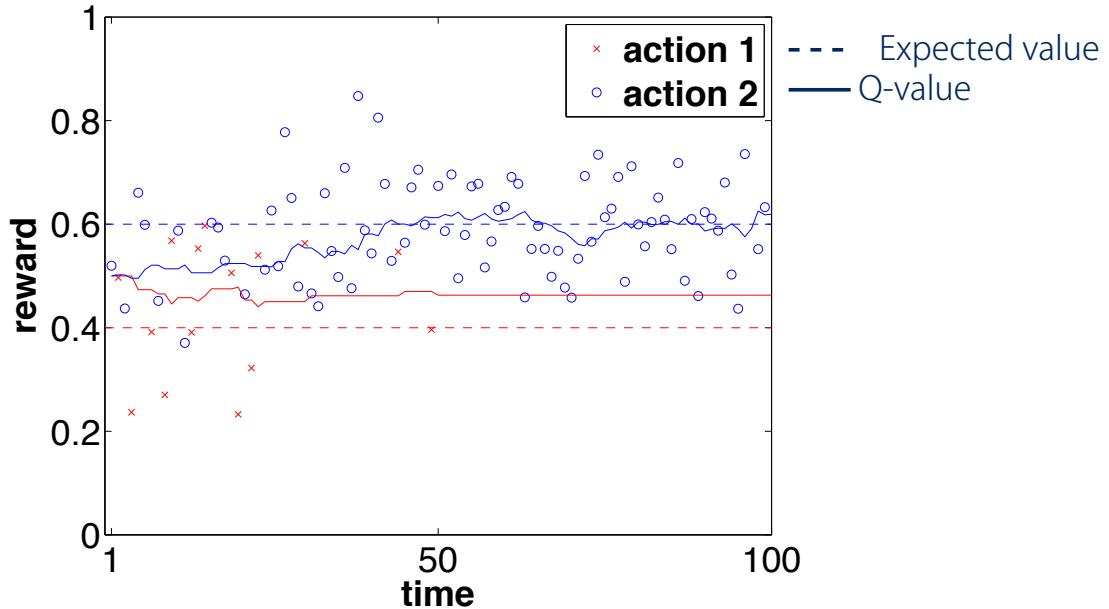
# Action values in dynamic environments

Static environment, off-policy Q-value updates



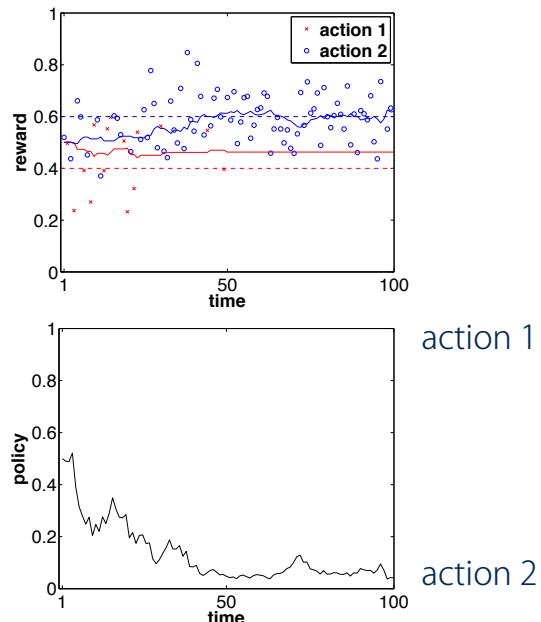
# Action values in dynamic environments

Static environment, on-policy Q-value updates



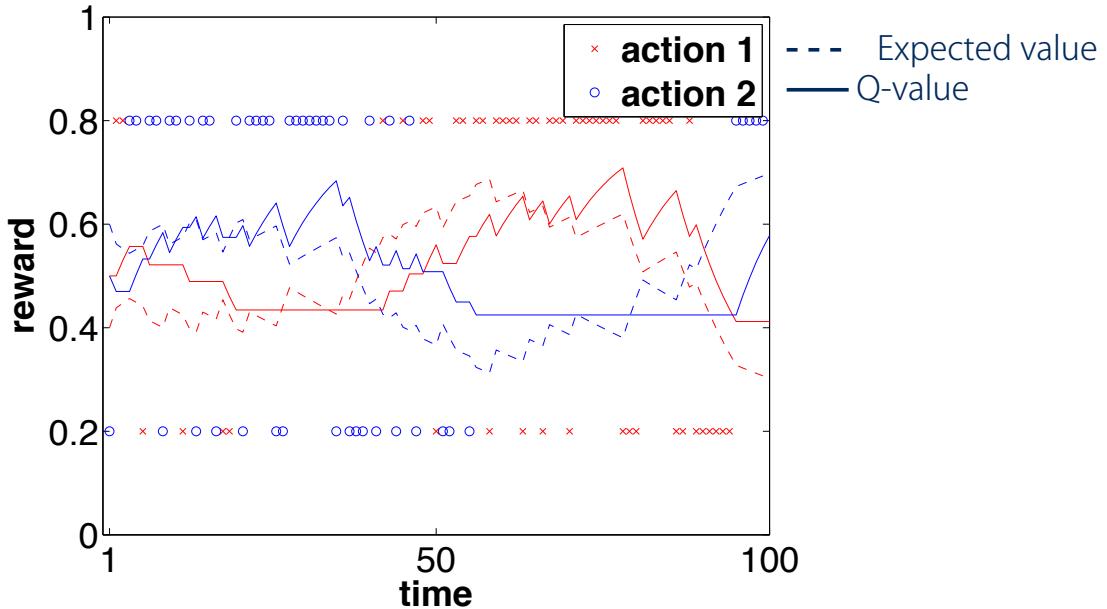
# Action values in dynamic environments

Static environment, on-policy Q-value updates



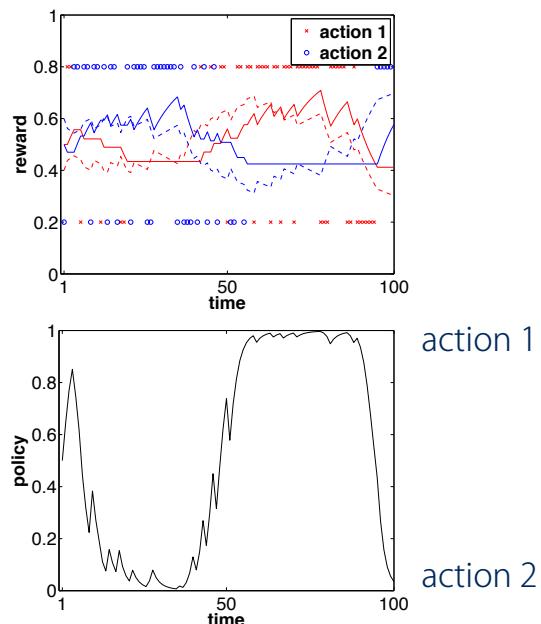
# Action values in dynamic environments

Adversarial environment, on-policy Q-value updates



# Action values in dynamic environments

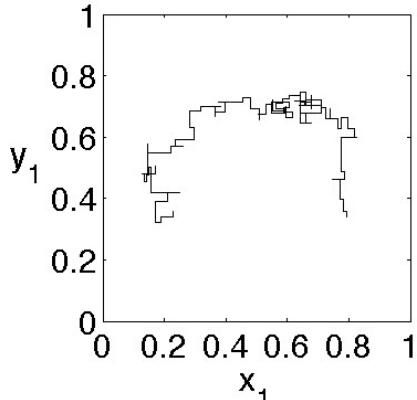
Adversarial environment, on-policy Q-value updates



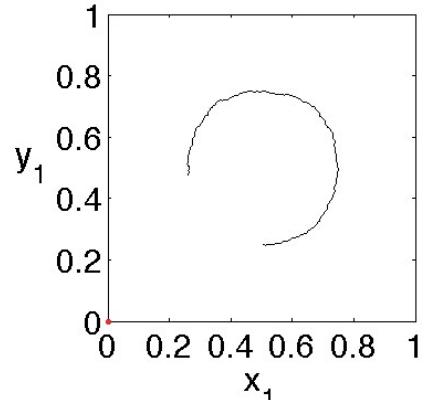
# Deriving Learning Dynamics

Matching Pennies

	$H$	$T$
$H$	1, 0	0, 1
$T$	0, 1	1, 0



$$\alpha = 0.1$$



$$\alpha = 0.001$$

# Deriving Learning Dynamics

Learning algorithm  $\lim_{\alpha \rightarrow 0} (E(\Delta x)) = \frac{dx}{dt} = \dot{x}$  Dynamical system

Advantages of dynamical systems

- ▶ Deterministic
- ▶ Convergence guarantees using Jacobian
- ▶ Vast related body of literature (e.g., bifurcation theory)

# Deriving Learning Dynamics

Example: Cross learning [Boerger97]

$$x_i(t+1) \leftarrow \begin{cases} (1 - \alpha r_i)x_i + \alpha r_i & \text{if } i \text{ selected} \\ (1 - \alpha r_j)x_i & \text{for other action } j \text{ selected} \end{cases}$$

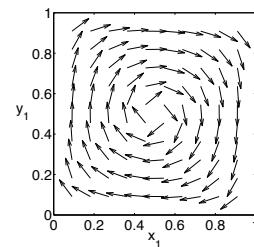
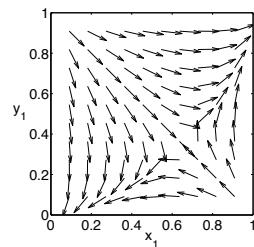
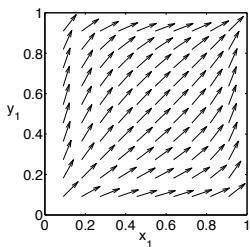
$$\begin{aligned} E(\Delta x_i) &= x_i [(1 - \alpha r_i)x_i + \alpha r_i - x_i] + \sum_{k \neq i}^n x_k [(1 - \alpha r_k)x_i - x_i] \\ &= \alpha x_i [(1 - x_i)r_i - \sum_{k \neq i}^n x_k r_k] = \alpha x_i [r_i - \sum_k^n x_k r_k] \end{aligned}$$

Learning algorithm  $\lim_{\alpha \rightarrow 0} (E(\Delta x)) = \frac{dx}{dt} = \dot{x}$  Dynamical system

$$\dot{x}_i = x_i \left[ E[r_i] - \sum_k^n x_k E[r_k] \right] = \underbrace{x_i [(Ay)_i - xAy]}_{\text{replicator dynamics}}$$

# Deriving Learning Dynamics

	Prisoners' Dilemma	Battle of Sexes	Matching Pennies	
	$D \quad C$	$B \quad S$	$H \quad T$	
$D$	$\begin{array}{ c c } \hline 1, 1 & 5, 0 \\ \hline 0, 5 & 3, 3 \\ \hline \end{array}$	$\begin{array}{ c c } \hline 2, 1 & 0, 0 \\ \hline 0, 0 & 1, 2 \\ \hline \end{array}$	$\begin{array}{ c c } \hline 1, -1 & -1, 1 \\ \hline -1, 1 & 1, -1 \\ \hline \end{array}$	$x_1$
$C$	$y_1 \quad 1 - y_1$	$y_1 \quad 1 - y_1$	$y_1 \quad 1 - y_1$	$1 - x_1$



$$\dot{x}_i = x_i [(Ay)_i - xAy]$$

$$\dot{y}_i = y_i [(xB)_i - xBy]$$

# Deriving Learning Dynamics

---

Dynamics have been derived for

- ▶ Learning Automata (Cross Learning)
- ▶ Regret Matching (RM)
- ▶ Variations of Infinitesimal Gradient Ascent
  - ▶ Infinitesimal Gradient Ascent (IGA)
  - ▶ Win-or-Learn-Fast (WoLF) IGA
  - ▶ Weighted Policy Learning (WPL)
- ▶ Variations of Q-learning
  - ▶ Repeated Update Q-learning  
*See our talk on Thursday, Session F4 Learning 1, 13:50*
  - ▶ Frequency Adjusted Q-learning

## Illustrating Convergence

---

Q-learning [Watkins92]

$x_i$  probability of playing action  $i$

$\alpha$  learning rate

$r$  reward

$\tau$  temperature

Update rule

$$Q_i(t+1) \leftarrow Q_i(t) + \alpha \left( r_i(t) + \gamma \max_j Q_j(t) - Q_i(t) \right)$$

Policy generation function

$$x_i(Q, \tau) = \frac{e^{\tau^{-1} Q_i}}{\sum_j e^{\tau^{-1} Q_j}}$$

# Illustrating Convergence

---

Frequency Adjusted Q-learning (FAQ-learning) [Kaisers2010]

$x_i$  probability of playing action  $i$

$\alpha$  learning rate

$r$  reward

$\tau$  temperature

Update rule

$$Q_i(t+1) \leftarrow Q_i(t) + \alpha \frac{1}{x_i} \left( r_i(t) + \gamma \max_j Q_j(t) - Q_i(t) \right)$$

Policy generation function

$$x_i(Q, \tau) = \frac{e^{\tau^{-1} Q_i}}{\sum_j e^{\tau^{-1} Q_j}}$$

# Illustrating Convergence

---

Cross Learning [Boergers97]

$$\dot{x}_i = x_i \left[ E[r_i(t)] - \sum_k^n x_k E[r_k(t)] \right]$$

Frequency Adjusted Q-learning [Tuyls05, Kaisers2010]

$$\dot{x}_i = \alpha x_i \left( \tau^{-1} \left[ E[r_i(t)] - \sum_k^n x_k E[r_k(t)] \right] - \log x_i + \sum_k x_k \log x_k \right)$$

Proof of convergence in two-player two-action games  
[Kaisers2011, Kianercy2012]

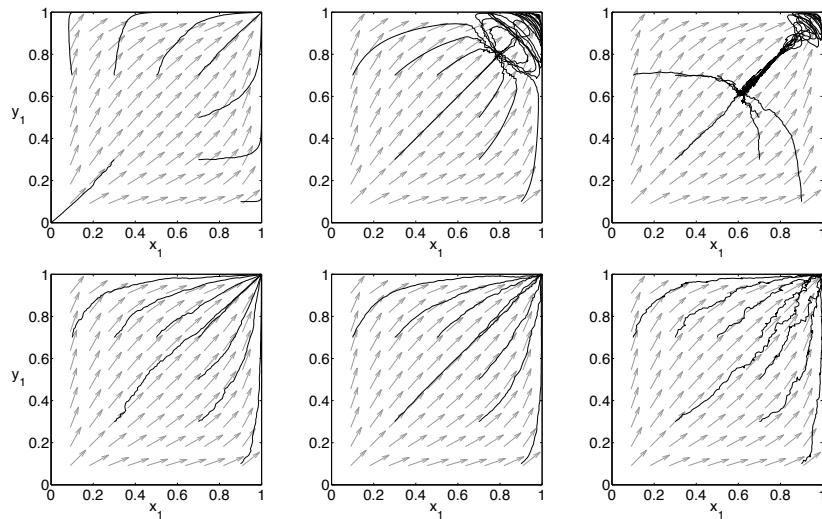
# Illustrating Convergence

Prisoners' Dilemma

Q

FAQ

$Q_{init}$



pessimistic

neutral

optimistic

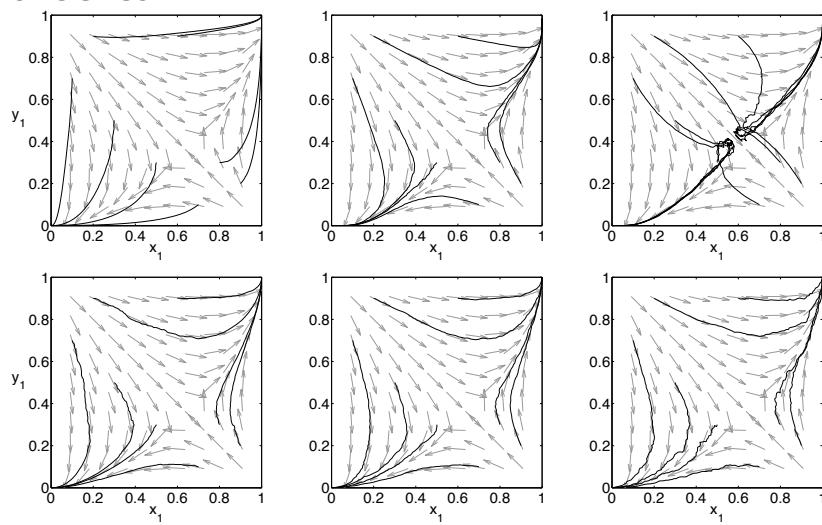
# Illustrating Convergence

Battle of Sexes

Q

FAQ

$Q_{init}$



pessimistic

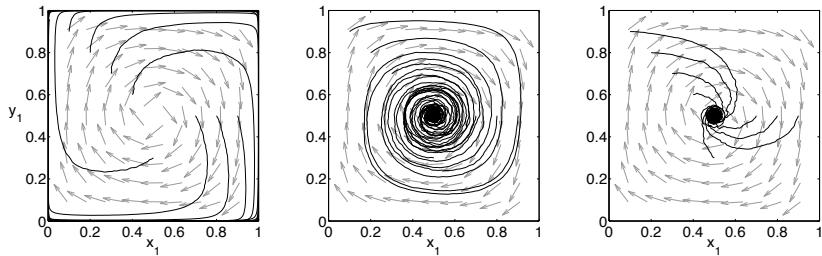
neutral

optimistic

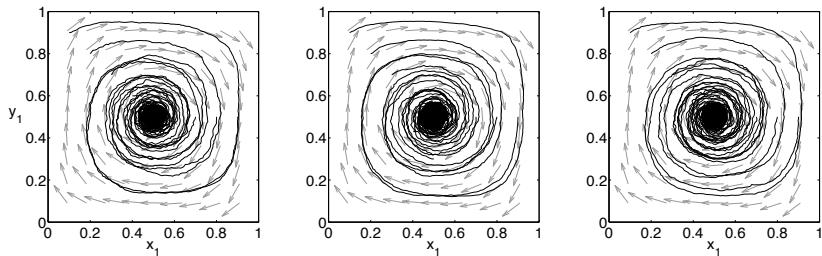
# Illustrating Convergence

## Matching Pennies

Q



FAQ



$Q_{init}$

pessimistic

neutral

optimistic

# Comparing Dynamics

Dynamical systems have been associated with

- ▶ Infinitesimal Gradient Ascent (IGA)
- ▶ Win-or-Learn-Fast Infinitesimal Gradient Ascent (WoLF)
- ▶ Weighted Policy Learning (WPL)
  
- ▶ Cross Learning (CL)
- ▶ Frequency Adjusted Q-learning (FAQ)
- ▶ Regret Matching (RM)

# Comparing Dynamics

Learning dynamics for two-agent two-action games. The common gradient is abbreviated  $\vec{\partial} = [yhAh^T + A_{12} - A_{22}]$ .

Algorithm	$\dot{x}$
IGA	$\alpha \vec{\partial}$
WoLF	$\vec{\partial} \cdot \begin{cases} \alpha_{min} & \text{if } V(x, y) > V(x^e, y) \\ \alpha_{max} & \text{otherwise} \end{cases}$
WPL	$\alpha \vec{\partial} \cdot \begin{cases} x & \text{if } \vec{\partial} < 0 \\ (1-x) & \text{otherwise} \end{cases}$
CL	$\alpha x(1-x) \vec{\partial}$
FAQ	$\alpha x(1-x) [\vec{\partial} \cdot \tau^{-1} - \log \frac{x}{1-x}]$
RM	$\alpha x(1-x) \vec{\partial} \cdot \begin{cases} (1+\alpha x \vec{\partial})^{-1} & \text{if } \vec{\partial} < 0 \\ (1-\alpha(1-x)\vec{\partial})^{-1} & \text{otherwise} \end{cases}$

# Comparing Dynamics

Cross Learning is linked to the replicator dynamics

$$\dot{x}_i = x_i \left[ E[f_i(t)] - \sum_k^n x_k E[f_k(t)] \right]$$

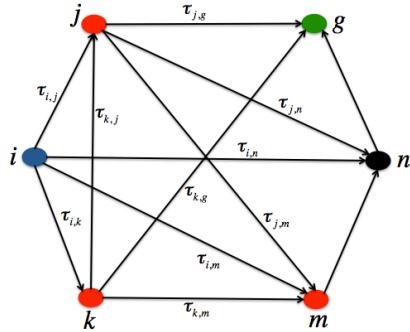
Gradient based algorithms use the orthogonal projection function, leading to the following dynamics for IGA:

$$\dot{x}_i = \alpha \left[ E[f_i(t)] - \sum_k^n \frac{1}{n} E[f_k(t)] \right]$$

Gradient dynamics use information about all actions, and are equivalent to the replicator dynamics under the uniform policy (i.e., also given off-policy updates).

# Replicator Dynamics

Path finding with Ant Colony Optimization



Pheromones  $\tau$  and travel cost heuristic  $\eta$  lead to probabilistic selection of state transition  $x_{i,j}$  from  $i$  to  $j$ :

$$x_{i,j} = \frac{\tau_{i,j}^\alpha \eta^{\beta}}{\sum_c \tau_{i,c}^\alpha \eta_{i,c}^\beta}, \text{ with } \alpha, \beta \text{ tuning parameters}$$

# Replicator Dynamics

Pheromone trail reinforcement in Ant Colony Optimization

$$\tau_{i,j}(t+1) = (1 - \rho)\tau_{i,j}(t) + \sum_{m=1}^M \delta_{i,j}(t, m),$$

where  $\rho$  denotes the pheromone evaporation rate,  $M$  is the number of ants and  $\delta_{i,j}(t, m) = Q \frac{n_{i,j}}{L}$  with  $Q$  being a constant,  $n_{i,j}$  being the number of times edge  $(i, j)$  has been visited.

$$\begin{aligned} \dot{x}_{i,j} &= \left( \frac{\tau_{i,j}^\alpha \eta^{\beta}}{\sum_c \tau_{i,c}^\alpha \eta_{i,c}^\beta} \right)' = \alpha x_{i,j} \frac{\dot{\tau}_{i,j}}{\tau_{i,j}} - \alpha x_{i,j} \sum_c \frac{\dot{\tau}_{i,c}}{\tau_{i,c}} x_{i,c} \\ &= \underbrace{\alpha x_{i,j} \left( \Theta_{i,j} - \sum_k x_{i,k} \Theta_{i,k} \right)}_{\text{replicator dynamics}}, \quad \Theta_{i,j} = \frac{\dot{\tau}_{i,j}}{\tau_{i,j}} \end{aligned}$$

# Replicator Dynamics

---

$$\dot{x}_i = x_i \left[ E[f_i(t)] - \sum_k^n x_k E[f_k(t)] \right]$$

Relative competitiveness as encoded by the replicator dynamics models

- ▶ the selection operator in evolutionary game theory
- ▶ pheromone trail reinforcement in swarm intelligence
- ▶ and exploitation in reinforcement learning dynamics.

## Summary

---

In strategic interactions

- ▶ the action values change over time
- ▶ the joint learning is a complex stochastic system
- ▶ dynamics can be captured in dynamical systems
  - ▶ proof of convergence
  - ▶ similarity of dynamics despite different implementations
  - ▶ link between learning, evolution and swarm intelligence
- ▶ different assumptions about observability give rise to a menagerie of algorithms to choose from

# Questions?

---



# Thank you!

Michael Kaisers | michaelkaisers@gmail.com

# Scaling Multi-agent Reinforcement Learning

Peter Vranckx

Joint work with: Y-M De Hauwere, A. Rodriguez, A. Nowe



1

## Sparse Interactions in Multi-agent reinforcement learning



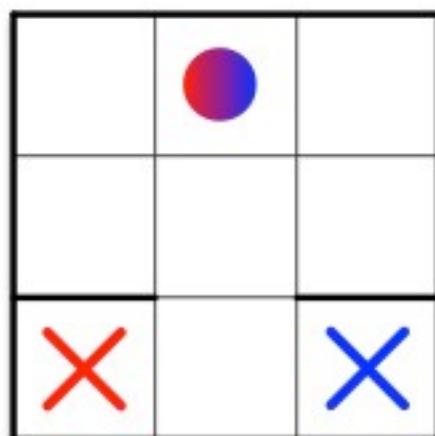
2

# Motivation

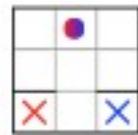
- Several issues arise when applying single agent RL techniques in multi-agent settings:
  - One vs. many learning agents?
  - Convergence? non-stationary, non-Markovian,...
  - Learning goal: e.g. maximize common reward vs. individual reward
  - Influence of action selection strategies and interactions
  - Credit assignment?
  - ...

3

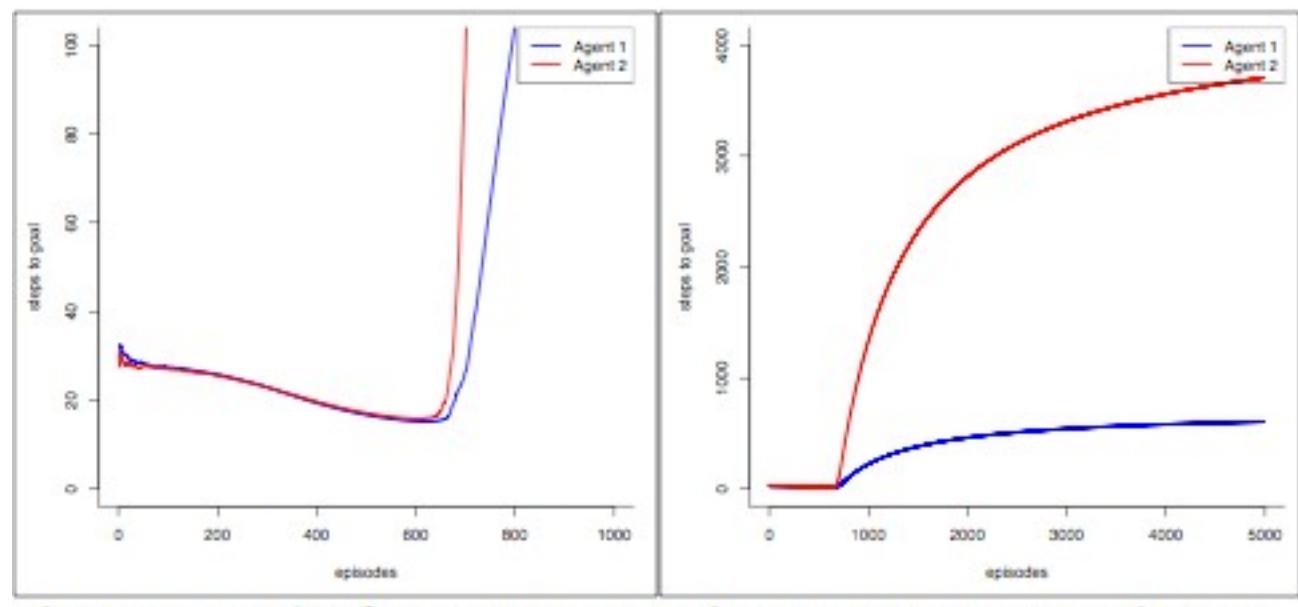
## Simple example



4



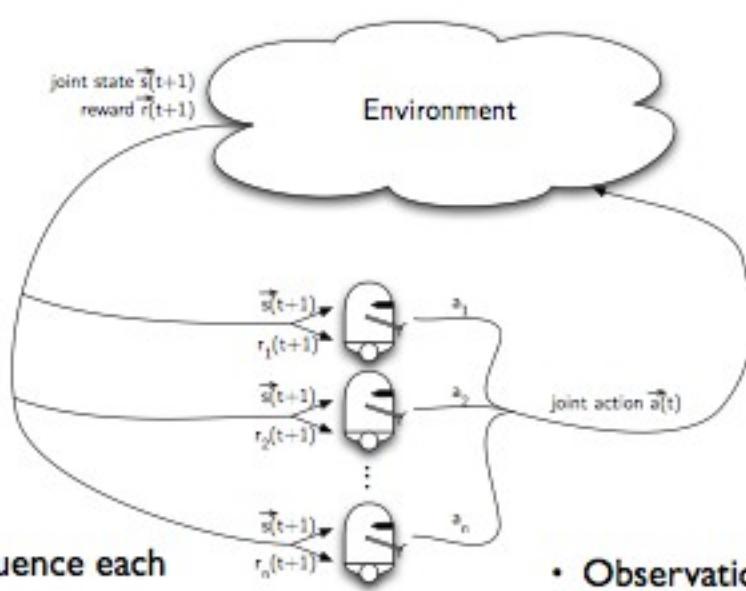
## Boltzmann exploration



Agents need information on other agents to coordinate

5

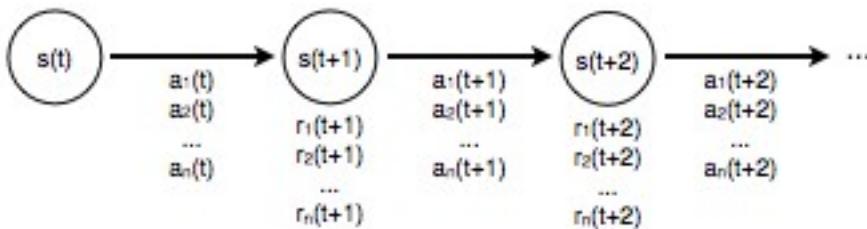
## Multi-agent reinforcement learning



- Agents influence each other
- Possibly conflicting interests
- Observations
- Expensive communication

6

# Markov Games



$n$

$$S = \{s^1, \dots, s^n\}$$

$$A = A_1 \times \dots \times A_n$$

$$T : S \times A_1 \times \dots \times A_n \times S \rightarrow [0, 1]$$

$$R_k : S \times A_1 \times \dots \times A_n \times S \rightarrow \mathbb{R}$$

- the number of agents
- a finite set of states
- with  $A_k$  the action set of agent k
- the transition function
- the reward function of agent k

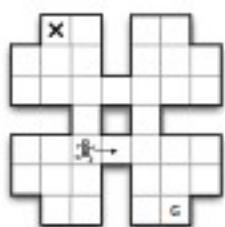
7

## Learning in Markov Games

- Learning occurs in joint state space (= all local information of all agents)
- Coordination mechanisms often require learning in joint action space
- Large information/communication requirements
- Exponential increases in problem size

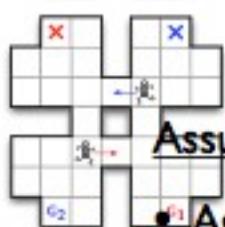
8

# Sparse interactions



## 1 agent

Transitions & rewards are only dependent on 1 agent

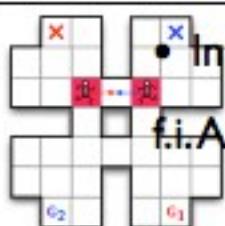


## 2 agents

Far away and not interacting with each other

Assumptions: Transitions & rewards are independent of state/

- Agents act and do their thing useful alone



- Interactions are sparse

## 2 agents

i.e. Air traffic control, automated warehouses, ...

Close to each other and interacting!!!

i.e. transitions & rewards are dependent

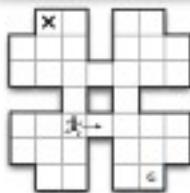
9

## Intuition of sparse interactions

Is there influence from another agent?

No

Act independently, as if single-agent.



Yes

Use a multi-agent technique to coordinate.



When should agents observe the state information of other agents to avoid coordination problems?

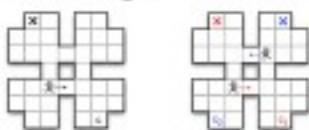
10

# Modeling interactions

- Dynamics of the system are a Markov game
- Model sparse interactions as a DEC-SIMDP (Melo et al., 2010)

$$\Gamma = (\underbrace{M^k}, \underbrace{(M^{I,l}, S^{I,l})})$$

MDP for each agent  $k$  in the absence of other agents (containing local states)



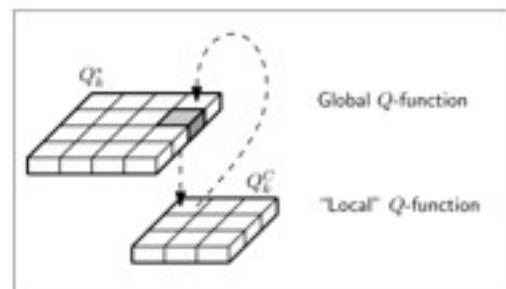
Team Markov game for the local interaction between  $K$  agents in  $L$  interaction states (containing system states)



11

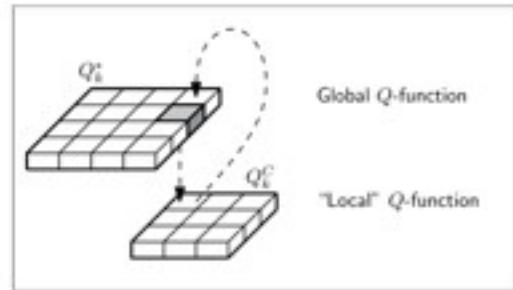
## Learning of Coordination

When to observe?



# Learning of Coordination

- Add Pseudo COORDINATE action
- External Active Perception
- Cost for coordination



13

## The algorithm

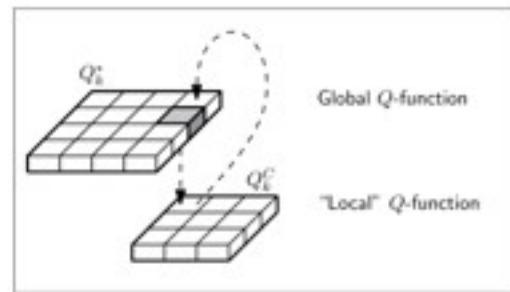
**Algorithm 1** Learning algorithm for agent  $k$

```
1: Initialize  $Q_k^*$  and  $Q_k^C$ ;  
2: Set  $t = 0$ ;  
3: while (FOREVER) do  
4:   Choose  $A_k(t)$  using  $\pi_a$ ;  
5:   if  $A_k(t) = \text{COORDINATE}$  then  
6:     if  $ActivePercept = \text{TRUE}$  then  
7:        $\hat{A}_k(t) = \pi_g(Q_k^C, X(t))$ ;  
8:     else  
9:        $\hat{A}_k(t) = \pi_g(Q_k^*, X_k(t))$ ;  
10:    end if  
11:    Sample  $R_k(t)$  and  $X_k(t+1)$ ;  
12:    if  $ActivePercept = \text{TRUE}$  then  
13:      QLUpdate( $Q_k^C; X(t), \hat{A}_k(t), R_k(t), X_k(t+1), Q_k^*$ );  
14:    end if  
15:   else  
16:     Sample  $R_k(t)$  and  $X_k(t+1)$ ;  
17:   end if  
18:   QLUpdate( $Q_k^*; X_k(t), A_k(t), R_k(t), X_k(t+1), Q_k^*$ );  
19:    $t = t + 1$ ;  
20: end while
```

14

# Utile Coordination

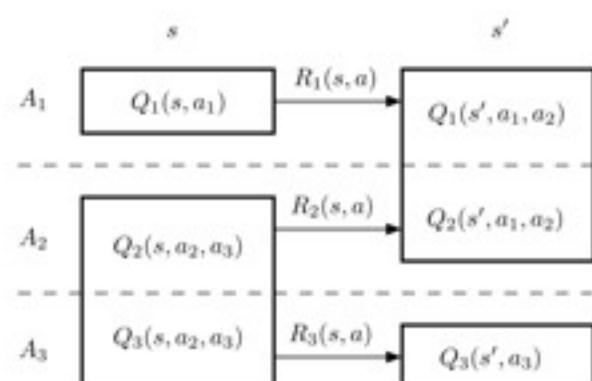
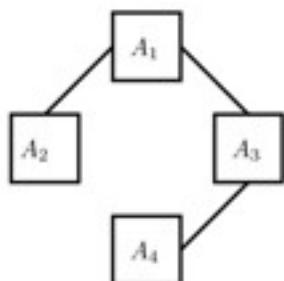
When to  
coordinate?



Kok & Vlassis, 2005

15

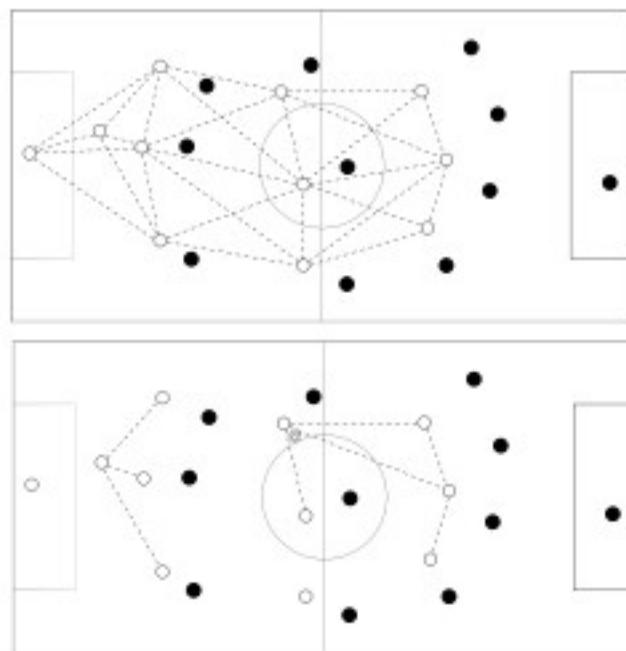
## Coordination graphs



Coordination through variable elimination  
algorithm

16

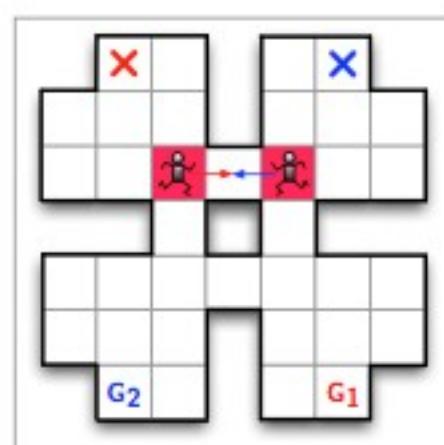
## Example: Robosoccer



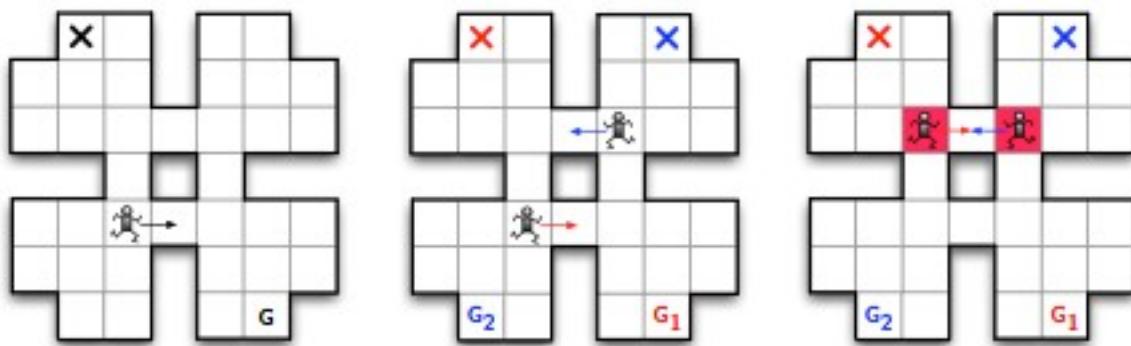
17

## CQ-Learning

Who to observe when?



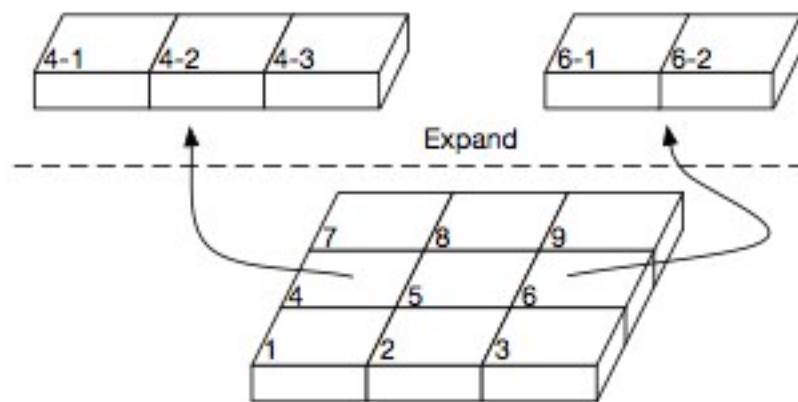
## Problem setting



- Agents only interact where their policies interfere
- Locally adapt policy

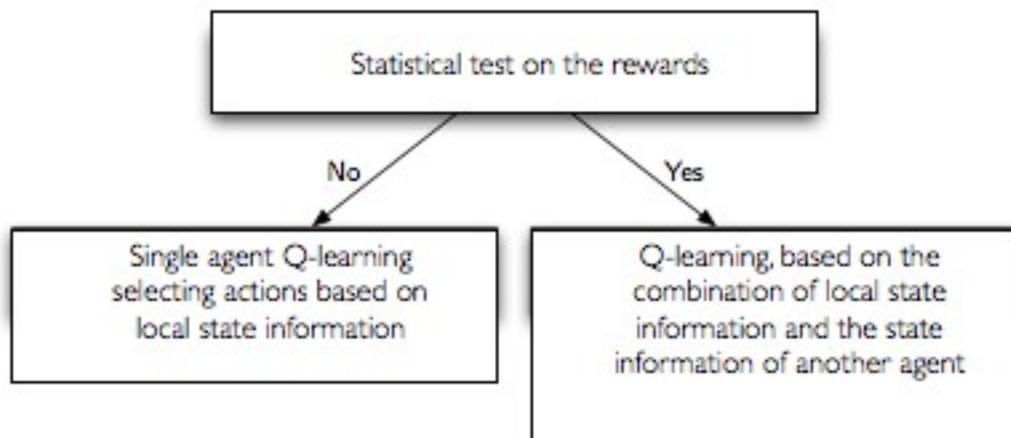
19

## Representation idea



20

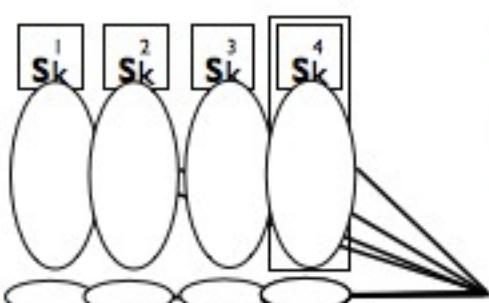
## Solution method: CQ-learning



21

## CQ-learninG

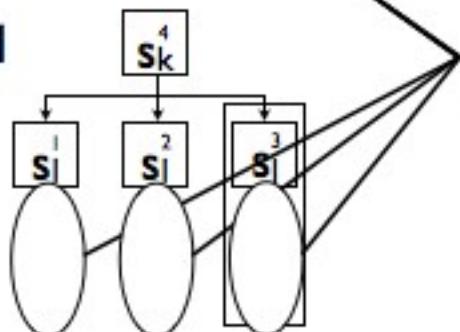
### STATISTICAL TESTS



Expected reward:

- Agents have been learning alone in the environment
- Agent  $k$  acts independently using only local state information ( $s_k$ ) in a multi-agent environment
- Performs statistical test against a baseline
- Samples its rewards, based on the state information of other agents & performs the same test  
 $s_k^4 \Rightarrow \langle s_k^4, s_l^3 \rangle$

Expand

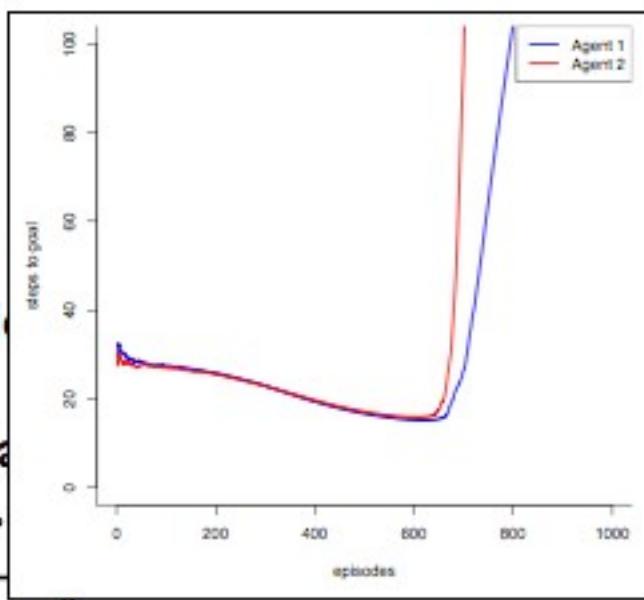
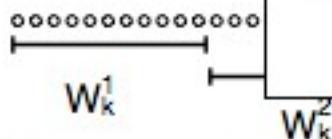


22

# CQ-LEARNING

BASELINE FOR STATISTICAL TESTS

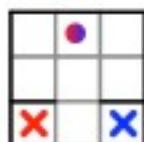
- Initial rewards (sliding window) for a particular state



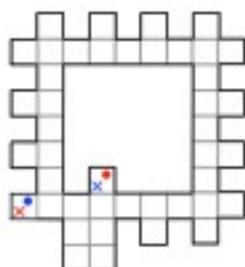
Compare  $W_k^1$  against  $W_k^2$

23

## Experimental results (1)



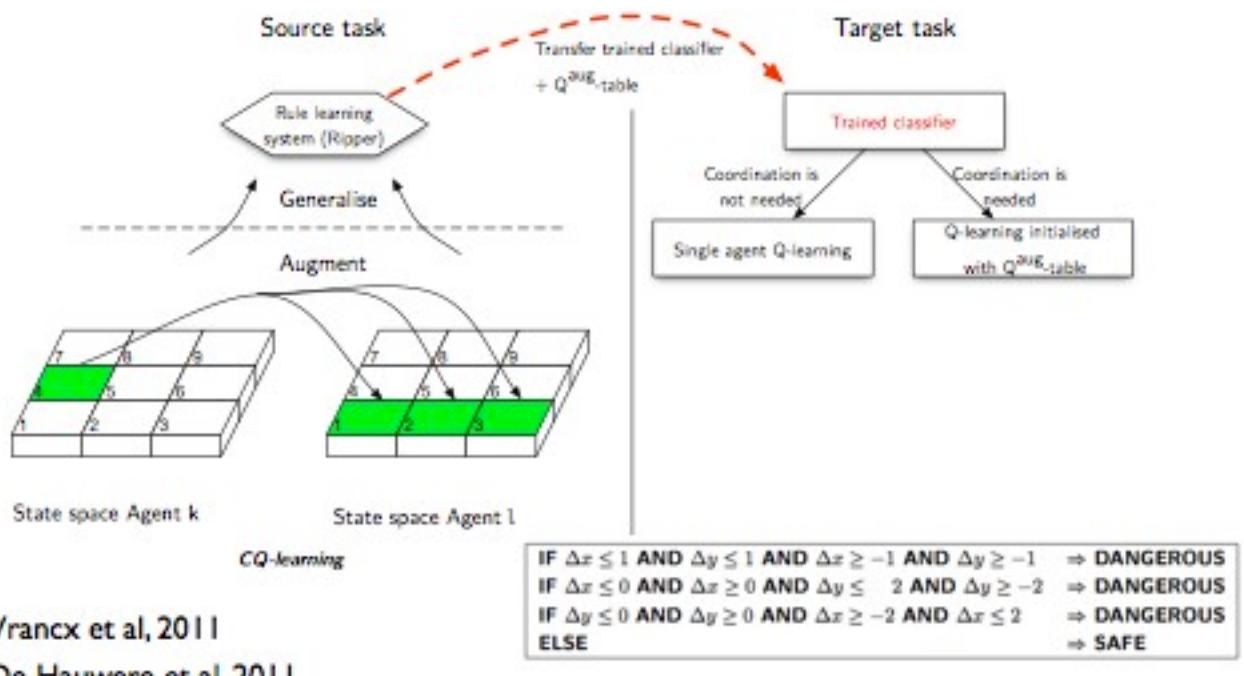
Env	Alg	#states	#actions	#coll	#steps
Grid-game-2 (min steps: 3)	Indep	9	4	2.7	$22.2 \pm 17.9$
	JS	81	4	0.1	$4.0 \pm 0.2$
	JSA	81	16	0.0	$4.7 \pm 0.1$
	LOC	$9.9 \pm 0.5$	5	0.1	$4.0 \pm 0.4$
	CQ	$10 \pm 0.0$	4	0.0	$3.6 \pm 0.3$
	CQ-NI	$10.9 \pm 2.0$	4	0.1	$4.0 \pm 0.3$



Env	Alg	#states	#actions	#coll	#steps
ISR (min steps: 4)	Indep	43	4	0.4	$9.3 \pm 44.8$
	JS	1849	4	0.1	$5.7 \pm 1.6$
	JSA	1849	16	0.0	$7.6 \pm 1.4$
	LOC	$51.3 \pm 82.3$	5	0.2	$6.7 \pm 7.5$
	CQ	$49.0 \pm 2.3$	4	0.1	$5.1 \pm 0.7$
	CQ-NI	$49.9 \pm 7.8$	4	0.1	$6.0 \pm 1.9$

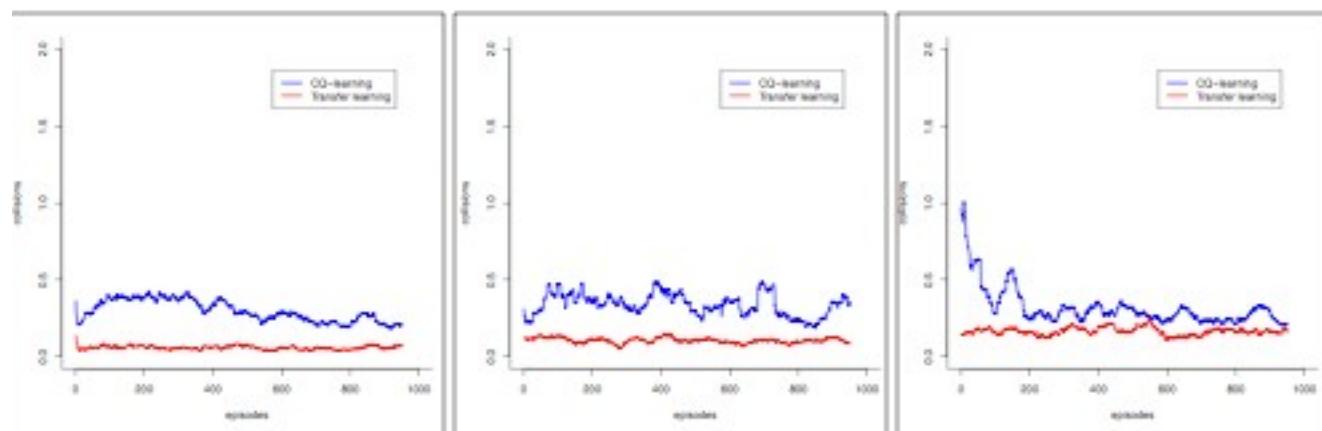
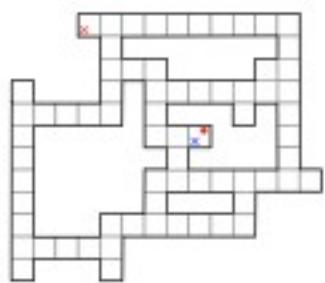
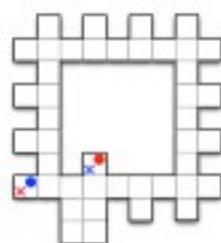
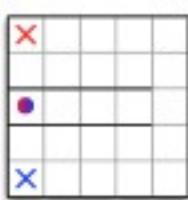
24

# Transfer learning with CQ-learning



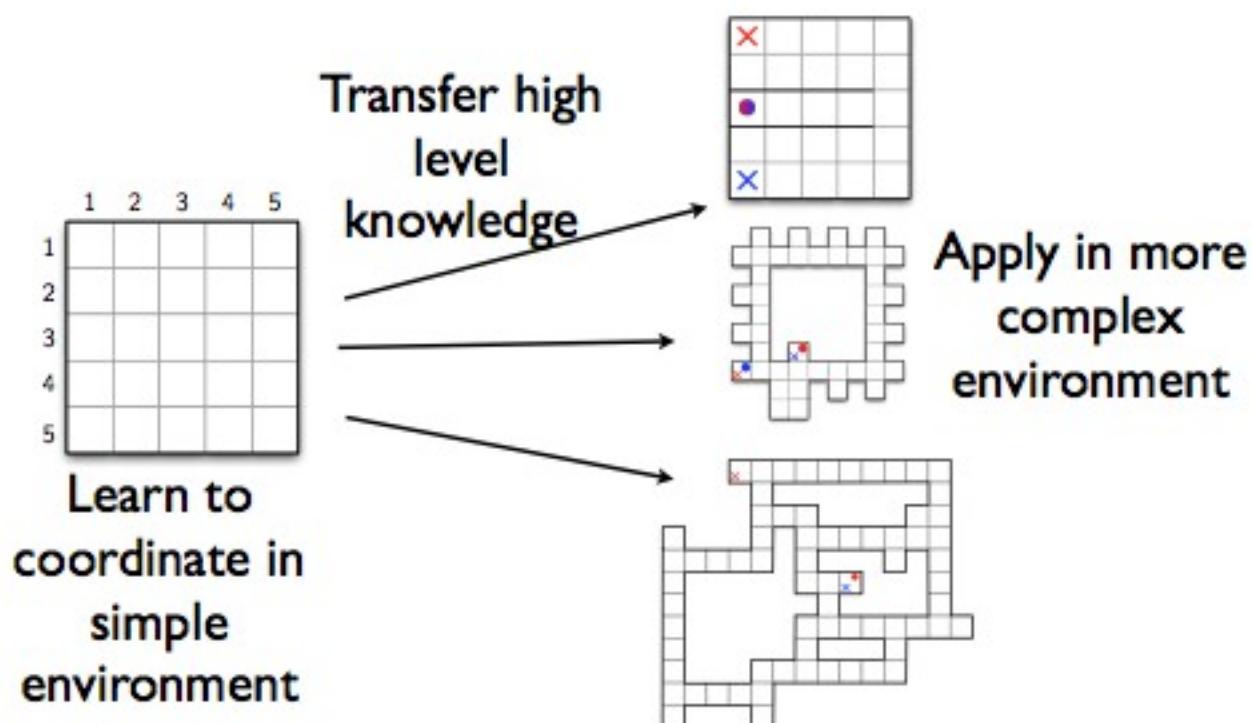
25

## Results (2)



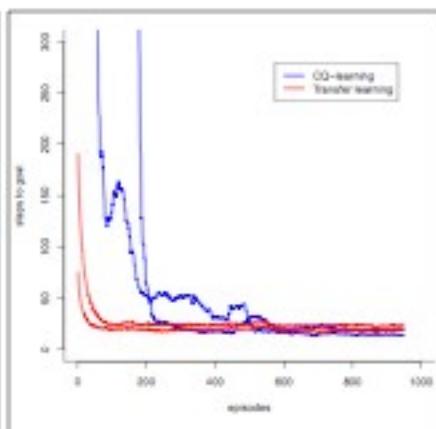
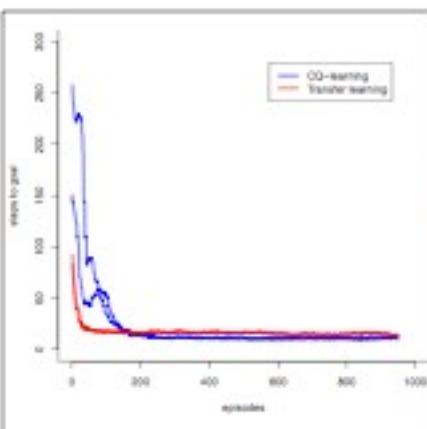
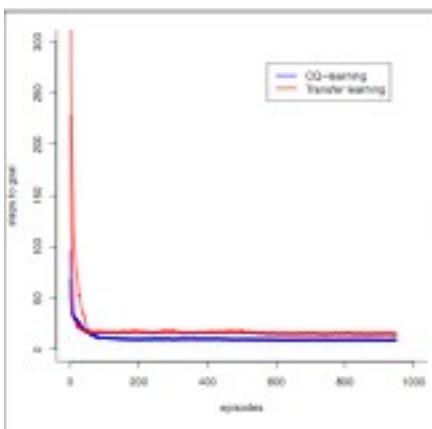
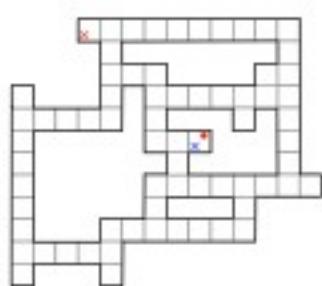
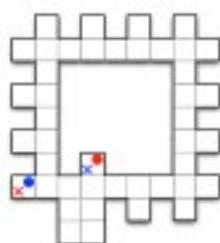
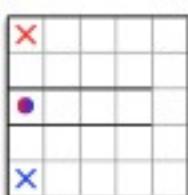
26

## Transfer learning with CQ-learning (2)



27

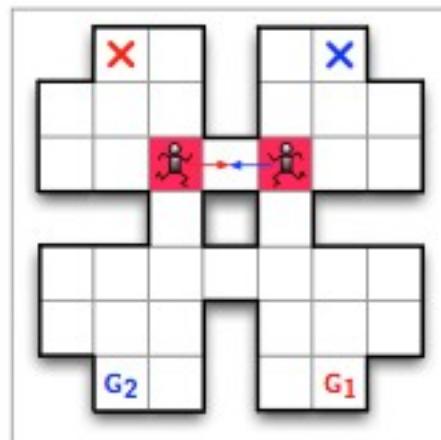
## Results



28

# FCQ-Learning

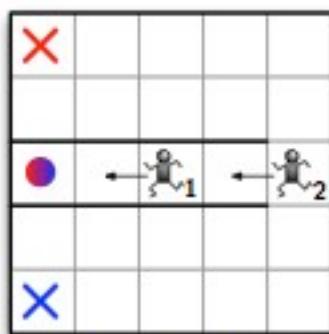
Who to observe when?



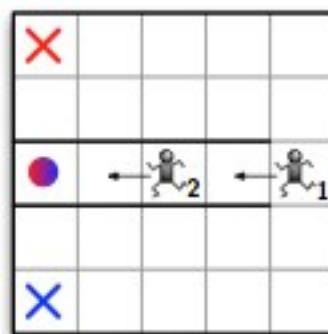
De Hauwere et al, 2011

29

## Problem setting



Reward: +20

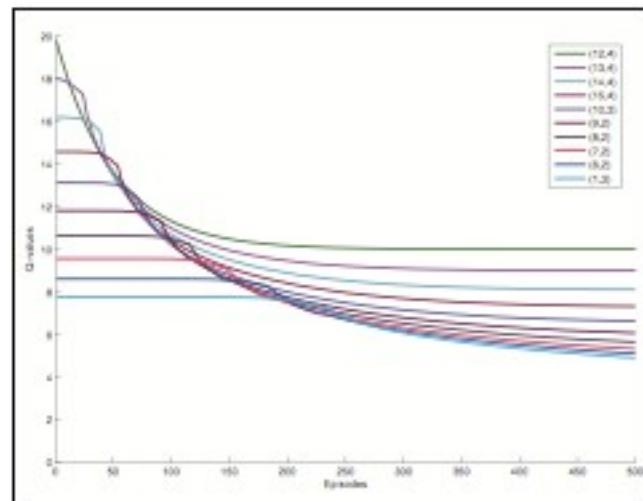
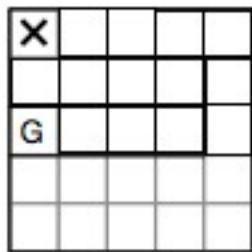


Reward: +10

- Reflected in immediate reward signal
- Too late to solve the problem

30

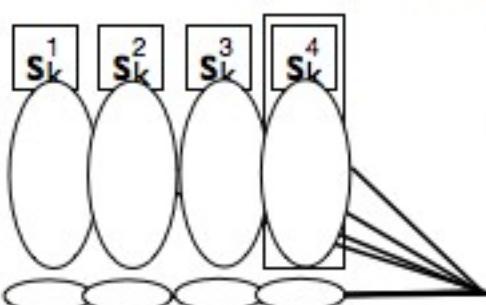
## Detecting relevant states



Changes in reward signal are reflected in the Q-values

31

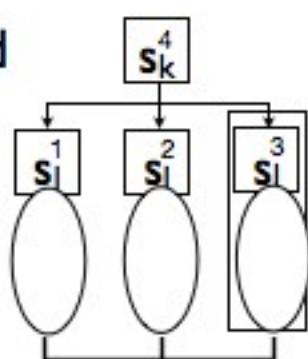
## FCQ-learning



statistical tests

- Agent k has been learning alone, and its Q-values have converged
- Agent k acts independently using only local state information ( $s_k$ ) in a multi-agent environment
- Performs statistical test against the single agent Q-values
- Samples rewards monte carlo and perform a comparison test to determine what information should be included

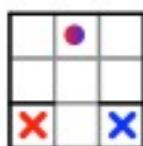
Expand



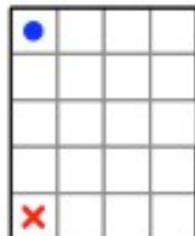
$$s_k^4 \Rightarrow \langle s_k^4, s_l^3 \rangle$$

32

## Experimental results



Environment	Algorithm	#states	#actions	#collisions	#steps	reward
Grid_game_2	Indep	9	4	2.4 ± 0.0	22.7 ± 30.4	-24.3 ± 35.6
	JS	81	4	0.1 ± 0.0	6.3 ± 0.3	18.2 ± 0.6
	LOC	9.0 ± 0.0	5	1.8 ± 0.0	10.3 ± 2.7	-6.8 ± 8.0
	FCQ	19.4 ± 4.4	4	0.1 ± 0.0	8.1 ± 13.9	17.6 ± 3.7
	FCQ.NI	21.7 ± 3.1	4	0.1 ± 0.0	7.1 ± 6.9	17.9 ± 0.7



Environment	Algorithm	#states	#actions	#collisions	#steps	reward
Bottleneck	Indep	43	4	n.a.	n.a.	n.a.
	JS	1849	4	0.0 ± 0.0	23.3 ± 30.8	13.1 ± 36.1
	LOC	54.0 ± 0.8	5	1.7 ± 0.6	167.2 ± 19,345.1	-157.5 ± 10,3
	FCQ	124.5 ± 32.8	4	0.1 ± 0.0	17.3 ± 1.3	16.6 ± 0.4
	FCQ.NI	135.0 ± 88.7	4	0.2 ± 0.0	19.2 ± 5.6	15.4 ± 2.3

33

## Conclusions

**In multi-agent environments with sparse interactions, learning these interaction states improves the learning process**

- Interaction states can be learned through increased penalties for miscoordination [Melo & Veloso, 2009]
- Interaction states can be identified using statistical tests on the reward signal (immediate + future) [De Hauwere et al, 2010 & 2011]
- Information about interaction states can be generalized and transferred between agents and environments [De Hauwere et al 2010, Vranckx et al 2010]

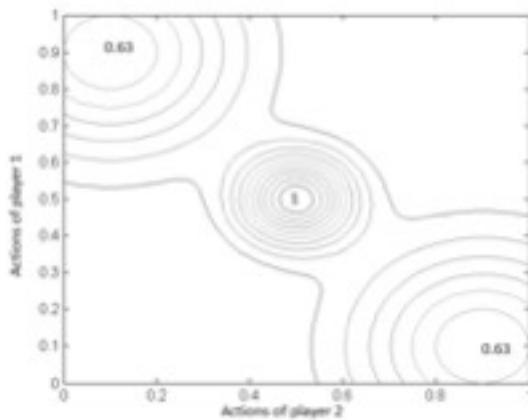
34

# Multi-agent Reinforcement Learning in Continuous Action Games



35

## Learning in Continuous Action Games

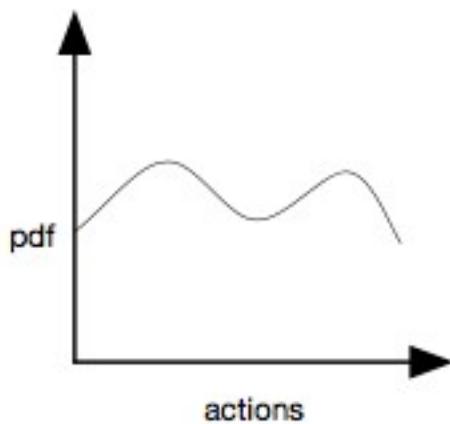


- Generalization of discrete normal form games
- Each agent now selects actions from continuous set
- Reward function is continuous function of all agents' actions

Assumptions: reward function is continuous,  
action sets are compact

36

# Continuous Action Reinforcement Learning Automata (CARLA)



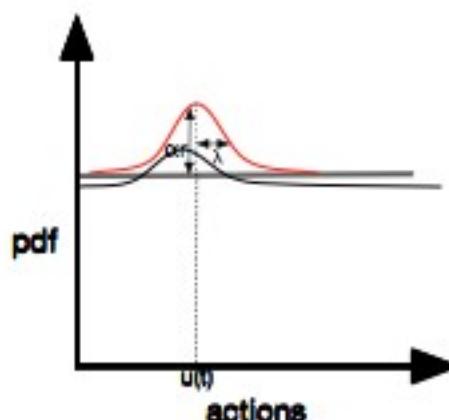
$$f_{k+1}(u) = \begin{cases} \eta_k \left( f_k(u) + \rho(u_k) \alpha e^{-\frac{1}{2} \left( \frac{u-u_k}{\lambda} \right)^2} \right) \\ 0 \end{cases}$$

- Extension of learning automata idea to continuous action spaces
- Now store continuous probability density distribution
- nonparametric pdf over actions
- parametric alternative: CALA (Santharam et al., 1994)

Howell, 1997

37

## CARLA Update



~~selected action, receive reinforcement, update corresponding distribution~~

- Extension of reward-inaction principles
- Reinforce selected action by adding Gaussian Bell to distribution
- Amplitude (strength of reinforcement) is determined by magnitude of reward and learning rate  $\alpha$
- Width (generalization) is determined by spreading rate parameter  $\lambda$

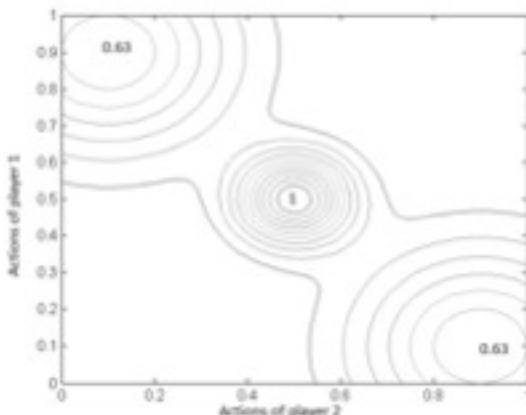
38

## CARLA Results

- In single agent systems: Converges to optimal neighborhood, depending on spreading rate  $\lambda$  (Rodriguez et al, 2011)
- In (cooperative) games a set of CARLA will converge to a locally superior strategy (Rodriguez et al, 2012)
- More accurate convergence can be achieved by transforming rewards and adaptively tuning the spreading rate  $\lambda$  (Rodriguez et al, 2011)

39

## Coordinated Exploration in Continuous Action Games

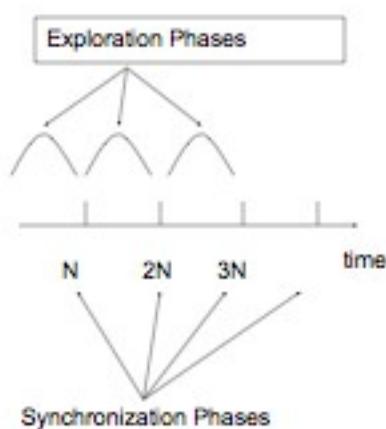


- In games CARLA may get stuck in local optima
- A narrow basin of attraction can make the global optimum difficult to find
- Coordinated exploration can allow learners to efficiently explore the joint action space

40

# Coordinated Exploration in discrete Games:

## Exploring selfish reinforcement learners (**ESRL**)



**Basic idea: 2 phases**

**Exploration: Be Selfish**

- Independent learning
- Convergence to different NE and Pareto Optimal non-NE

**Synchronization: Coordinate**

Verbeeck, 2004

41

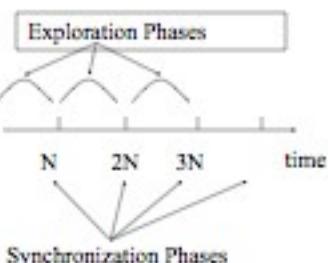
## ESRL in discrete games

The Penalty Game

Player B

10,10	0,0	k,k
0,0	2,2	0,0
k,k	0,0	10,10

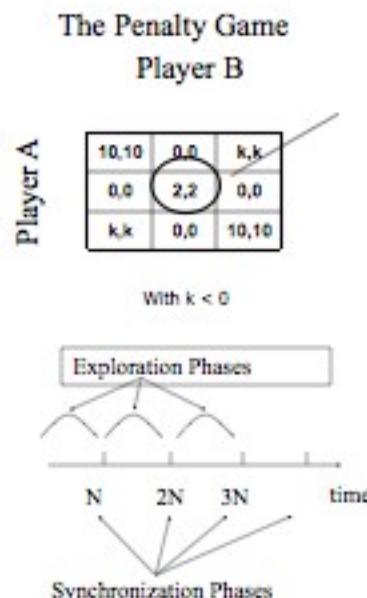
With  $k < 0$



42

# ESRL

- Exploration



- Use LRI → the agents converge to pure (Nash) joint action

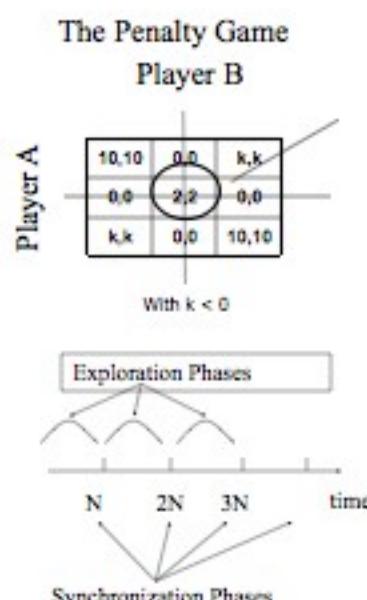
- Synchronization

- Update average payoff for action  $a$  converged to, optimistically
- Exclude action  $a$  and explore again if empty action set → RESET
- If done, select BEST

43

# ESRL

- Exploration



- Use LRI → the agents converge to pure (Nash) joint action

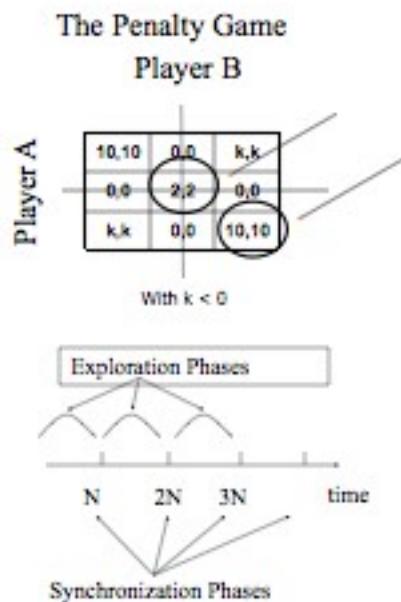
- Synchronization

- Update average payoff for action  $a$  converged to, optimistically
- Exclude action  $a$  and explore again if empty action set → RESET
- If done, select BEST

44

# ESRL

- Exploration



- Use LRI → the agents converge to pure (Nash) joint action

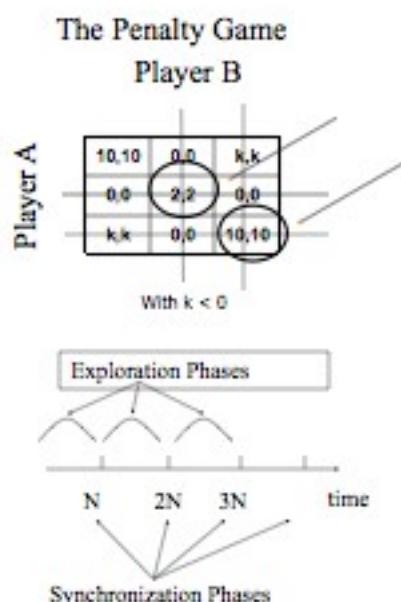
- Synchronization

- Update average payoff for action  $a$  converged to, optimistically
- Exclude action  $a$  and explore again if empty action set → RESET
- If done, select BEST

45

# ESRL

- Exploration



- Use LRI → the agents converge to pure (Nash) joint action

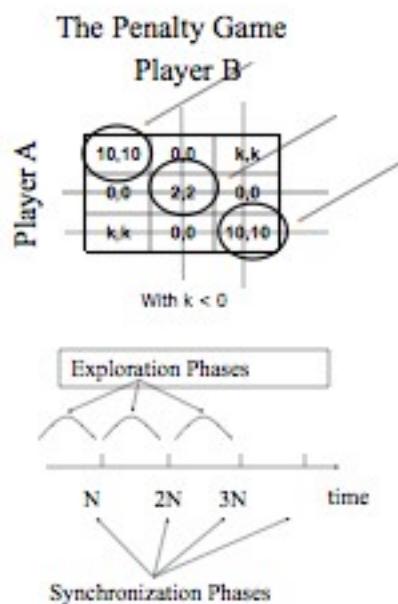
- Synchronization

- Update average payoff for action  $a$  converged to, optimistically
- Exclude action  $a$  and explore again if empty action set → RESET
- If done, select BEST

46

# ESRL

- Exploration



- Use LRI → the agents converge to pure (Nash) joint action

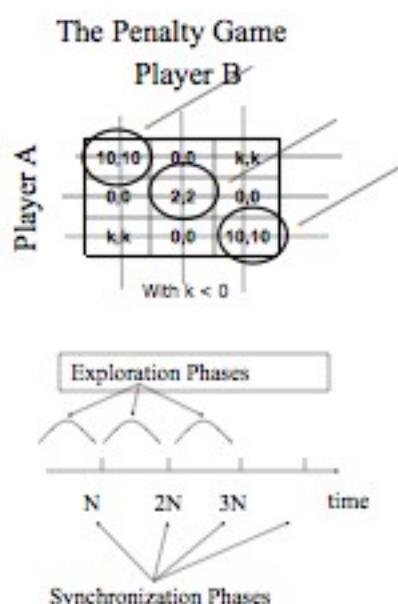
- Synchronization

- Update average payoff for action  $a$  converged to, optimistically
- Exclude action  $a$  and explore again if empty action set → RESET
- If done, select BEST

47

# ESRL

- Exploration



- Use LRI → the agents converge to pure (Nash) joint action

- Synchronization

- Update average payoff for action  $a$  converged to, optimistically
- Exclude action  $a$  and explore again if empty action set → RESET
- If done, select BEST

48

# ESRL

- Exploration

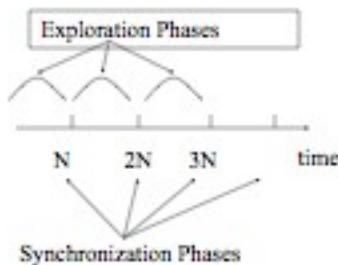
The Penalty Game

Player B

Player A

10,10	0,0	k,k
0,0	2,2	0,0
k,k	0,0	10,10

With  $k < 0$



- Use LRI → the agents converge to pure (Nash) joint action

- Synchronization

- Update average payoff for action  $a$  converged to, optimistically
- Exclude action  $a$  and explore again if empty action set → RESET
- If done, select BEST

49

## ESRL in conflicting interest games

### Battle of the sexes

Player 2

Player 1

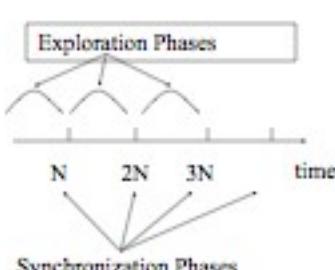
	B	S
B	2,1	0,0
S	0,0	1,2

- Exploration

- Use LRI → the agents converge to pure (Nash) joint action

- Synchronization

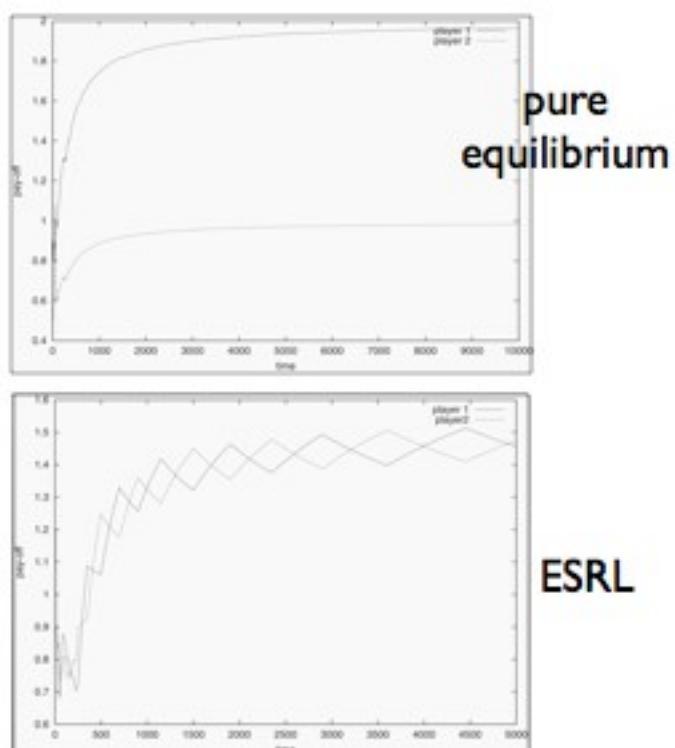
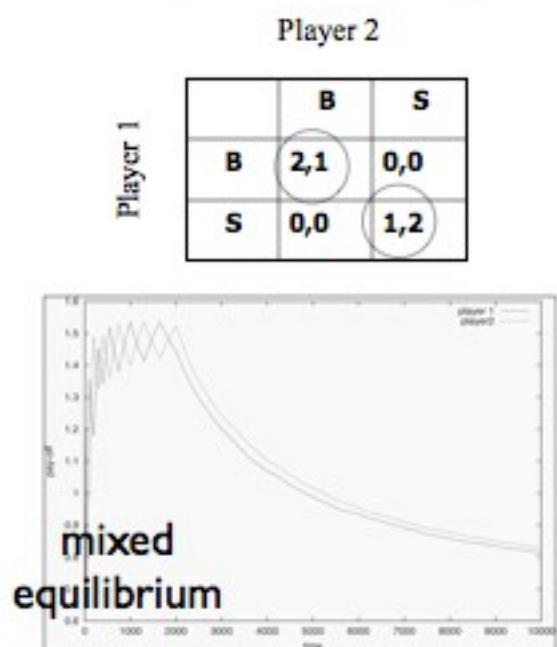
- Update average payoff for action  $a$  converged to, optimistically
- Exclude action  $a$  and explore again if empty action set → RESET
- Keep alternating to ensure fair payoffs



50

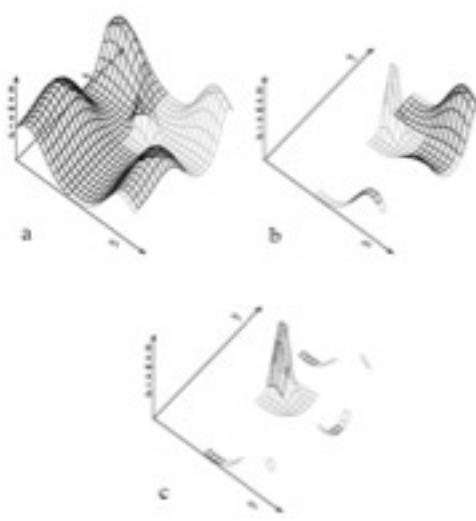
# ESRL in conflicting interest games

## Battle of the sexes



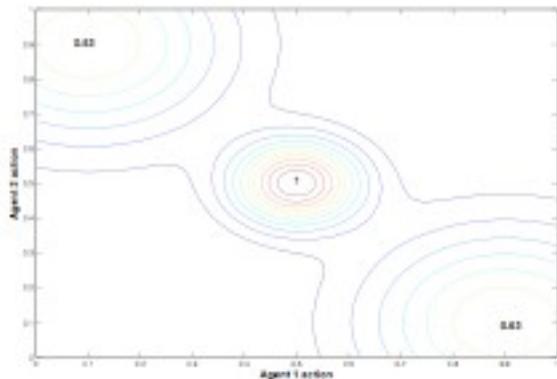
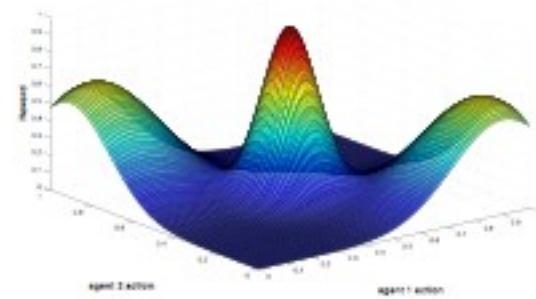
51

## ESRL in Continuous Settings



- Apply the same exploration / synchronization idea
- In continuous action spaces it makes no sense to exclude a single action
- We identify basin of attraction for learning outcomes, and eliminate entire region from the action space

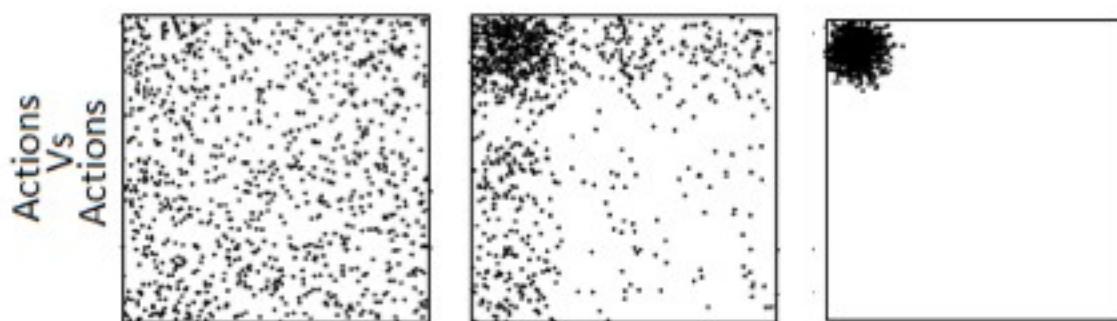
## Example Game



- 2 player continuous action game
- 3 local optima
- global optimum has smallest basin of attraction

53

## Identifying the Basin of Attraction



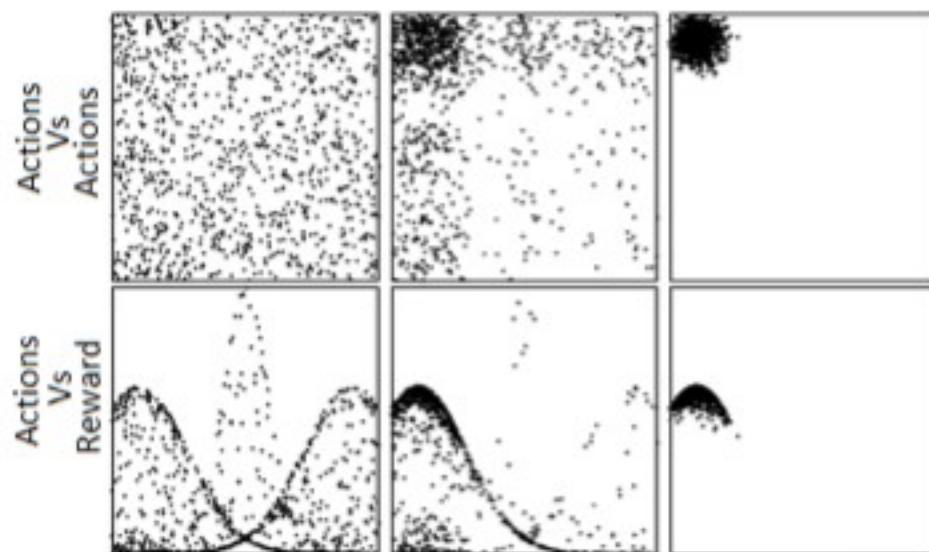
Initially learners explore randomly

Learners enter basin of attraction

Learners converge to attractor

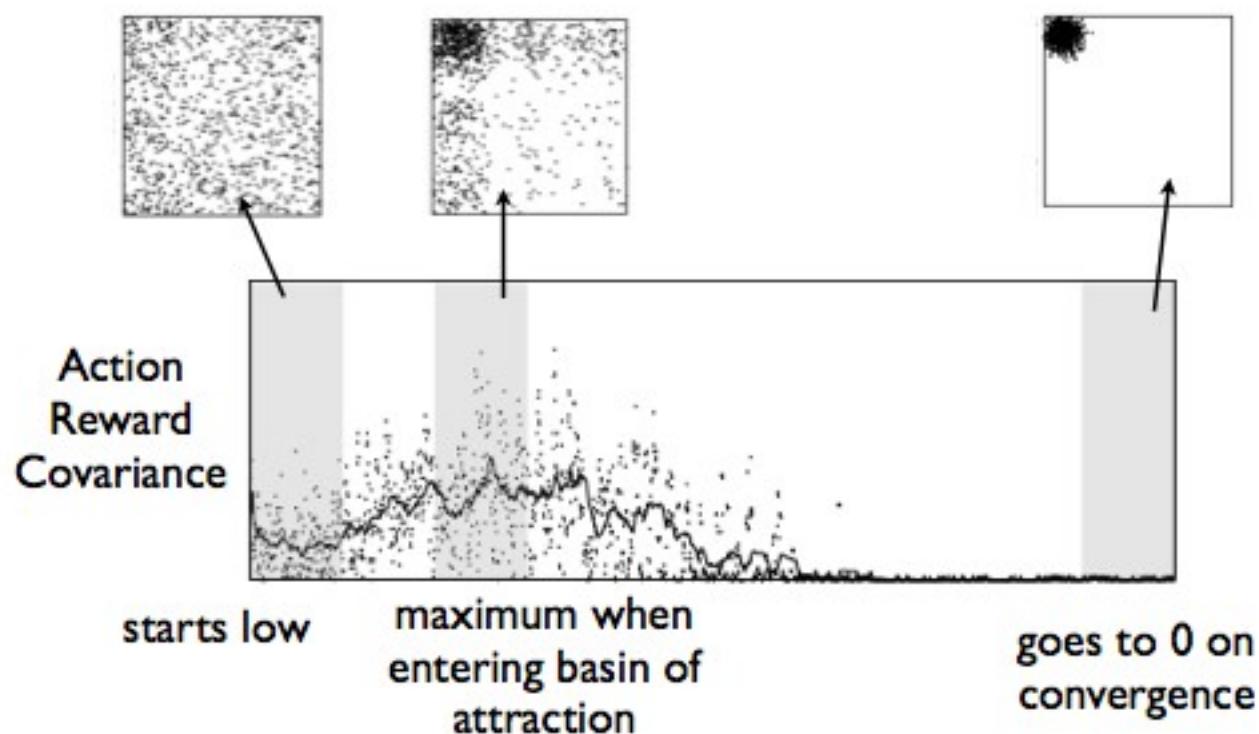
54

## Rewards During Learning



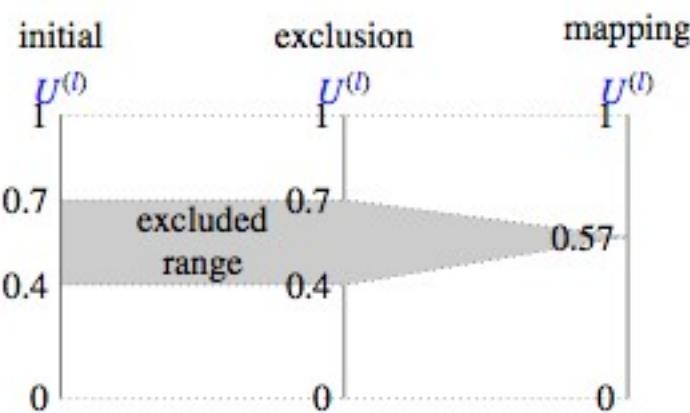
55

## Action - Reward Covariance



56

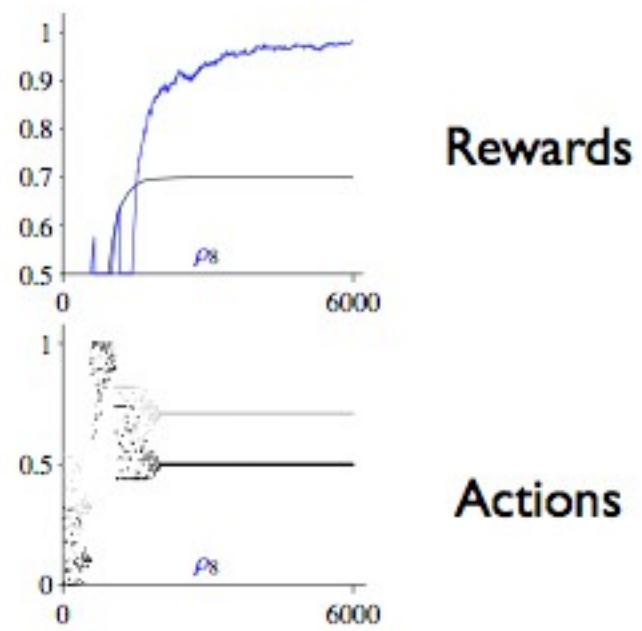
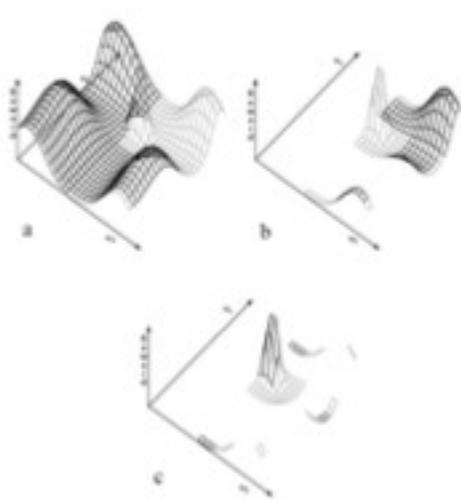
## Eliminate Action Range



- Use Covariance to identify action range
- Delete range from CARLA PDF
- Renormalize PDF

57

## Results



58

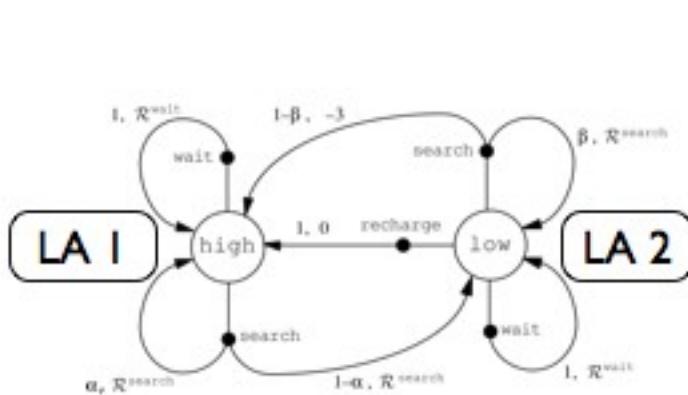
## Conclusions

- Learning automata formalisms can also be used in continuous action games
- Discrete coordination mechanisms can be extended to continuous case

59

## Multi-agent Reinforcement Learning in Continuous State Spaces

# Interconnected Learning Automata (ILA)



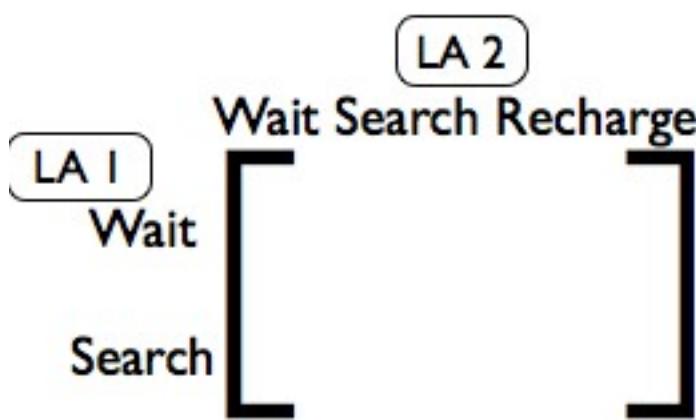
- Actor-Critic architecture
- LA is assigned to each state
- LA is updated using estimate of accumulated reward (state value) under current policy

Witten, 1977

Wheeler & Narendra, 1986

61

## ILA Analysis

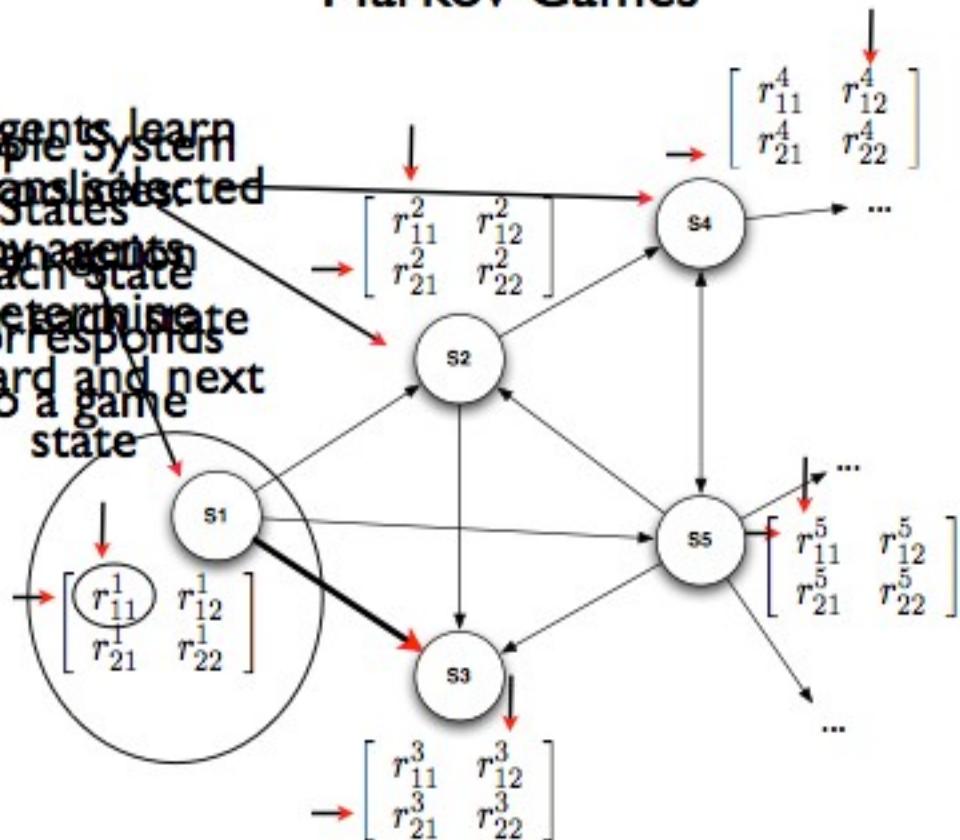


- Behavior of automata can be approximated by a game
- Equilibria of game represent optimal policies
- Game does not need to be explicitly calculated

62

## Markov Games

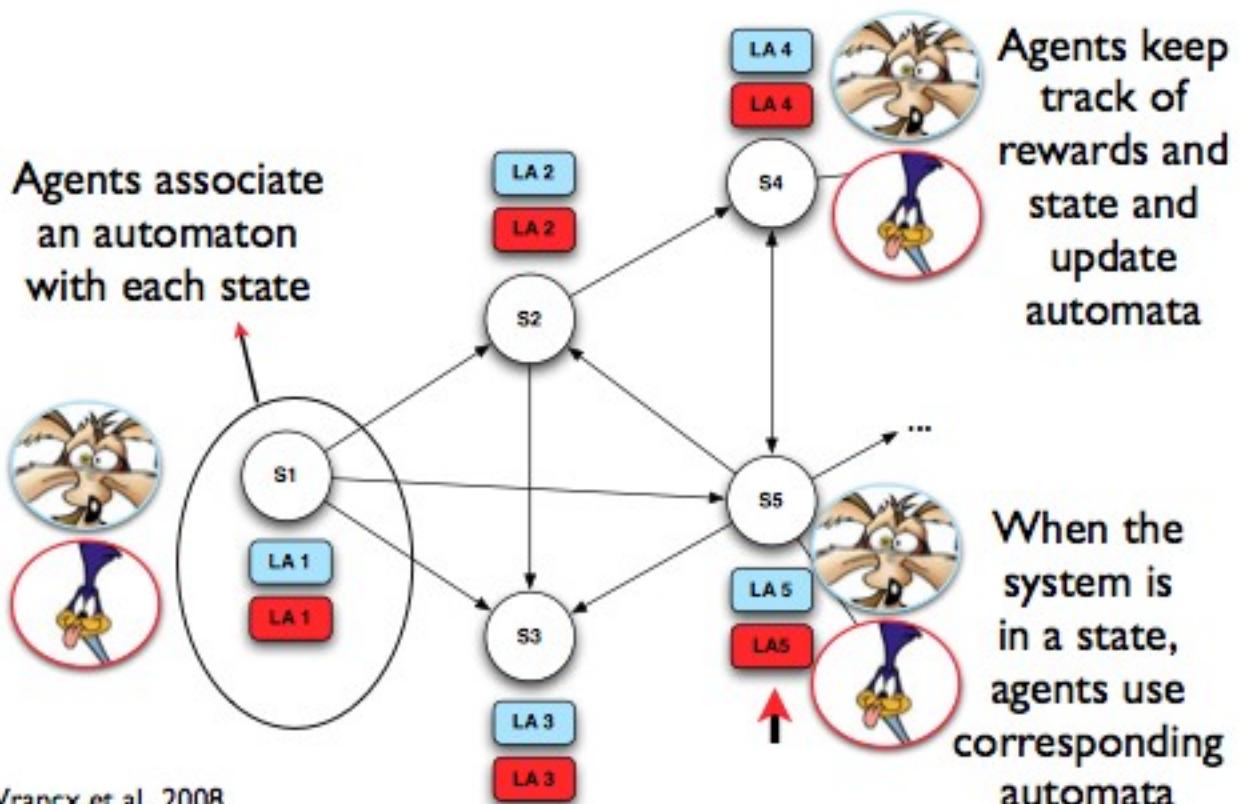
Agents learn  
Multiple System  
Actions selected  
States  
by agents  
Each state  
determine  
Corresponds  
to a game  
state



63

## MG-ILA Algorithm

Agents associate  
an automaton  
with each state

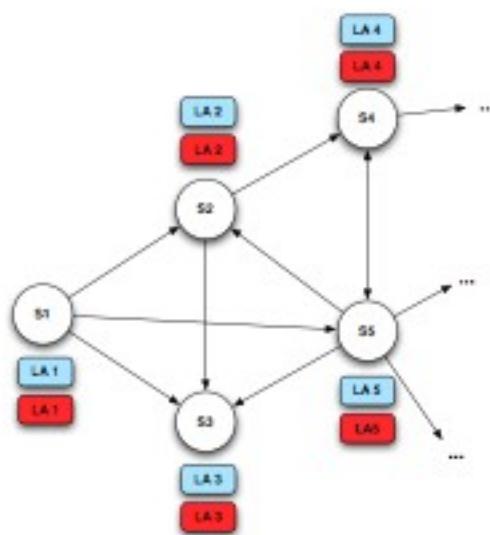


Agents keep  
track of  
rewards and  
state and  
update  
automata

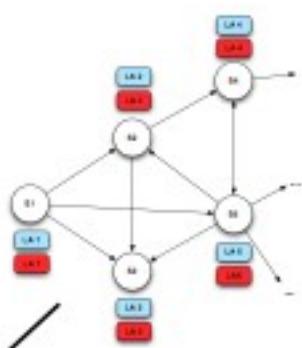
When the  
system is  
in a state,  
agents use  
corresponding  
automata

## MG-ILA Analysis

- Approximate learning behaviour by 2 games
- High level and low level view
- High level: agent interactions (Markov game)
- Low level: automata interactions
- Link both views



65



pol1 pol2 ... pol M

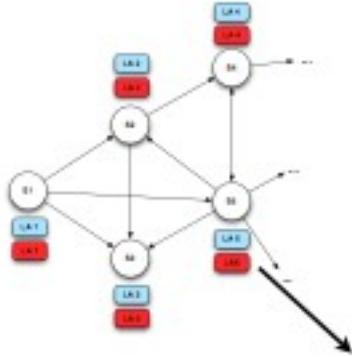


pol1  
pol2  
...  
pol M

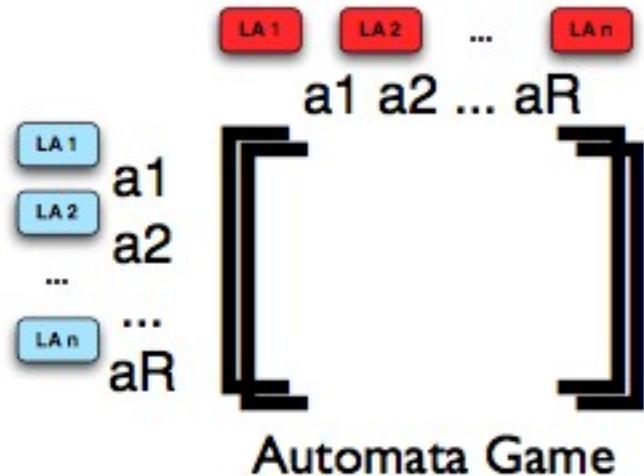
Agent Game

- Agents are players
- Look at agent Policies
- Payoff are expected rewards in Markov game
- High level view of outcome

66

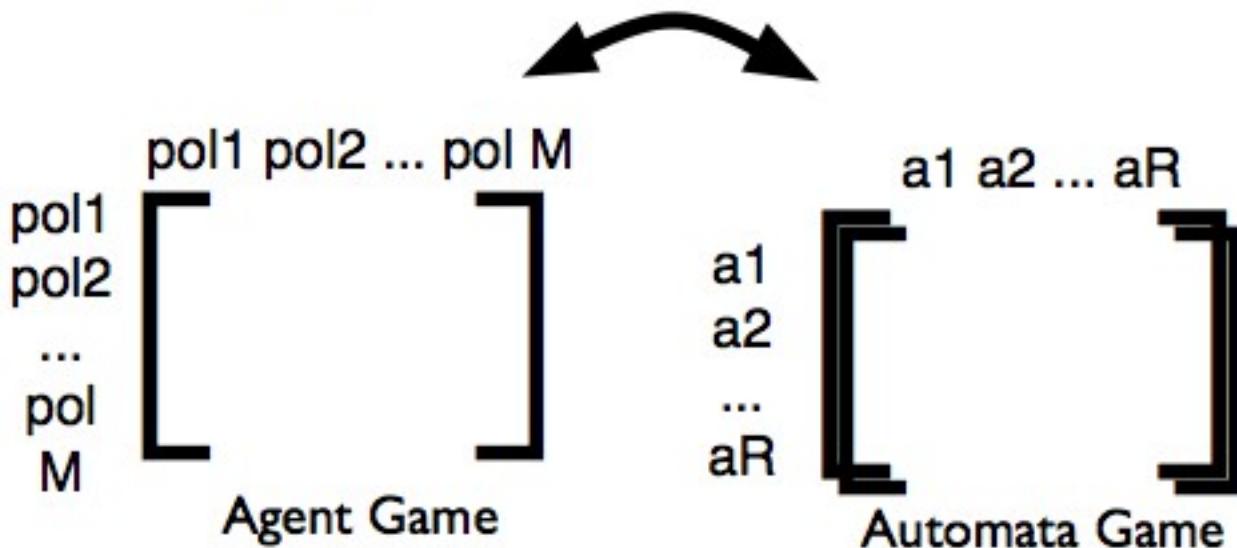


- Automata are players
- Each player selects action for 1 agent, 1 state
- Payoff are expected rewards in Markov game
- Low level view of interactions



67

Agents using automata find equilibrium between policies      Equilibria in both games correspond      Learning automata find equilibrium in automata game



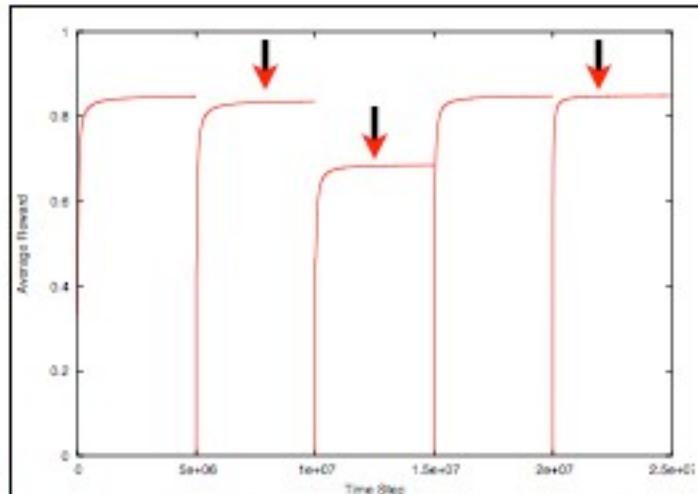
68

# Coordinated Exploration

Coordination mechanisms like ESRL can be used to find global optimum / achieve fair payoffs

0.85	0.76	0.61	0.18
0.48	0.75	0.30	0.33
0.82	<b>0.84</b>	0.27	0.32
0.43	0.90	<b>0.69</b>	0.44
0.31	0.33	0.36	0.50
0.32	0.33	0.51	<b>0.54</b>

Sub-optimal Equilibria can occur

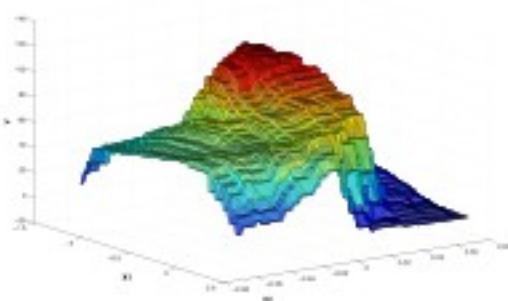


algorithm may converge to sub-optimal points

Vrancx et al., 2007

69

# Learning in Continuous MDPs



- State and Action space are continuous
- Exact tabular representations are no longer possible
- Approximation is needed to represent policies and value functions
- Approximate TD-algorithms (Q-learning/ SARSA) typically discretize action set to find greedy policy

70

# Linear Approximation

Value function approximation:

$$V(x) = \phi(x)^T \theta_v$$

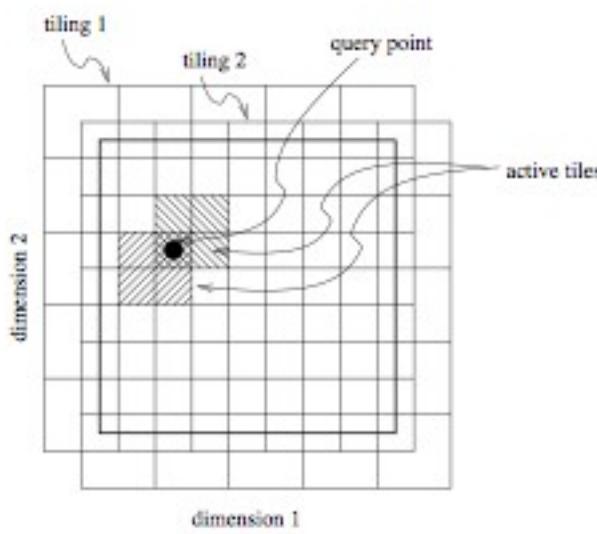
Policy approximation:

$$\pi(x) = \phi(x)^T \theta_u$$

- $\phi(x)$  are basis functions / features of state  $x$
- $\theta$  are learnt parameters
- approximation is linear in state features
- $\phi(x)$  can be non-linear functions of state variables

71

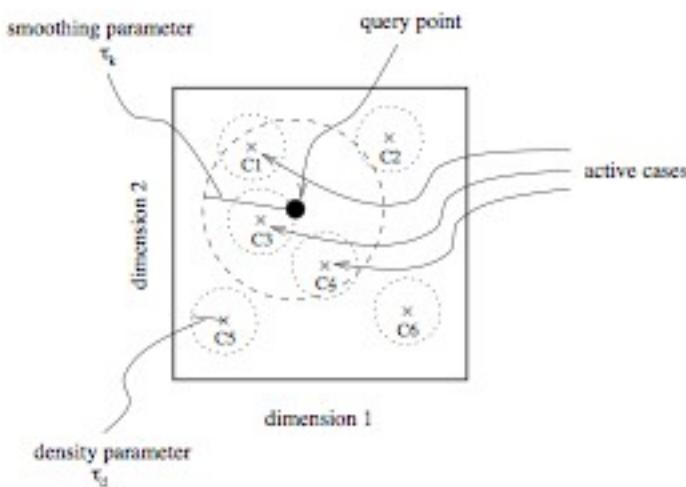
# Tile-Coding



- State space is covered with overlapping grids
- $\phi(x)$  is binary vector with 1 value for each tile
- In each grid only the tile in which state  $x$  falls is active
- $\phi(x)$  is 1 for active tiles, 0 else

72

## Kernel Based Approximation

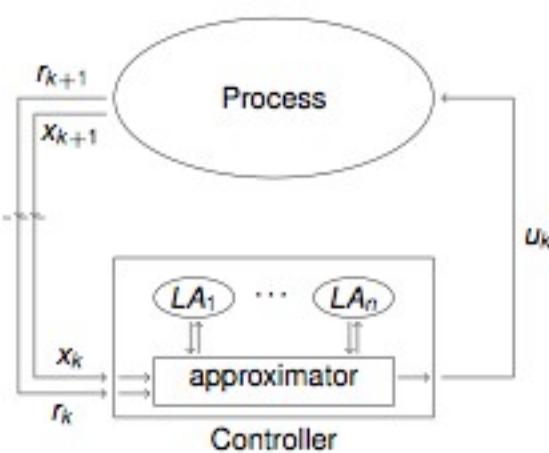


$$K(d) = e^{(d^2/\tau_k^2)}$$

- Select sample states  $C_1, C_2, \dots, C_n$
- For state  $x$ : select nearest neighbors based on distance  $d(x, C_i)$
- $\Phi(x)[i]$  is  $K(d(x, C_i))$  if  $d(x, C_i) < \tau_k, 0$  else.
- $K$  is kernel function (typically Gaussian)
- Can also be used instance based (add centers as needed)

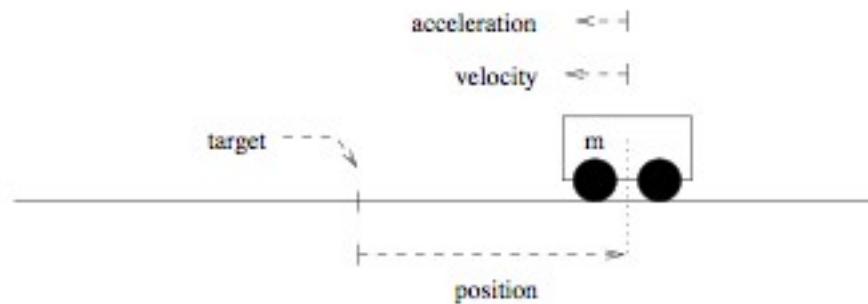
73

## CARLA learning in Continuous MDPs



- Same idea as ILA algorithm
- Set of CARLA learn policy (actor)
- Critic learns value of current policy
- Approximator is necessary to represent policy & values
- Continuous Action Game between CARLA with values as rewards

# Double Integrator



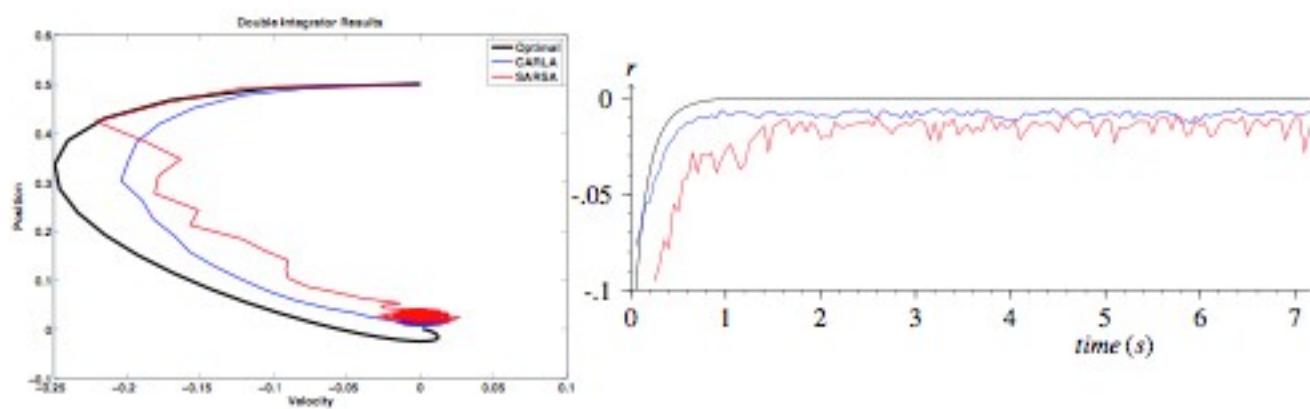
$$\mathbf{x}_t = [v_t; p_t]$$

$$u, v, p \in [-1, 1]$$

$$\mathbf{x}_{t+1} = \begin{bmatrix} v_{t+1} \\ p_{t+1} \end{bmatrix} = \begin{bmatrix} v_t \\ p_t \end{bmatrix} + \begin{bmatrix} 0 & 0 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} v_t \\ p_t \end{bmatrix} \Delta t + \begin{bmatrix} 1 \\ 0 \end{bmatrix} [a] \Delta t = \mathbf{x}_t + A\mathbf{x}_t \Delta t + B u_t \Delta t$$

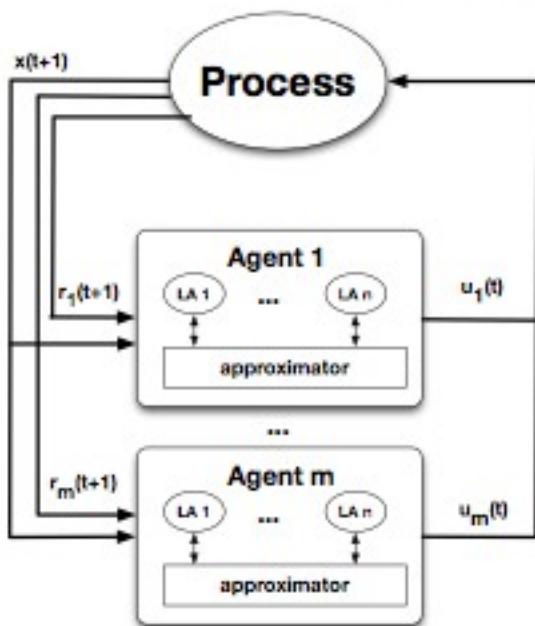
75

## Results: CARLA + tile coding



76

# Continuous Markov Games

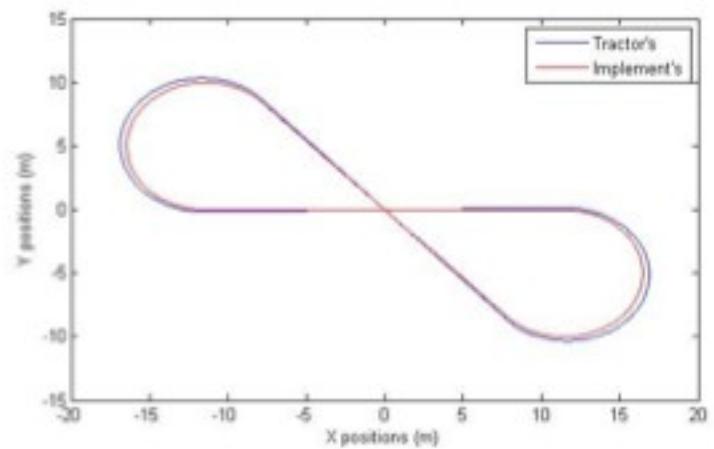


- Idea can also be extended to **Continuous Markov Games**
- Each agent now uses set of CARLA + approximator

Rodriguez et al., 2013b

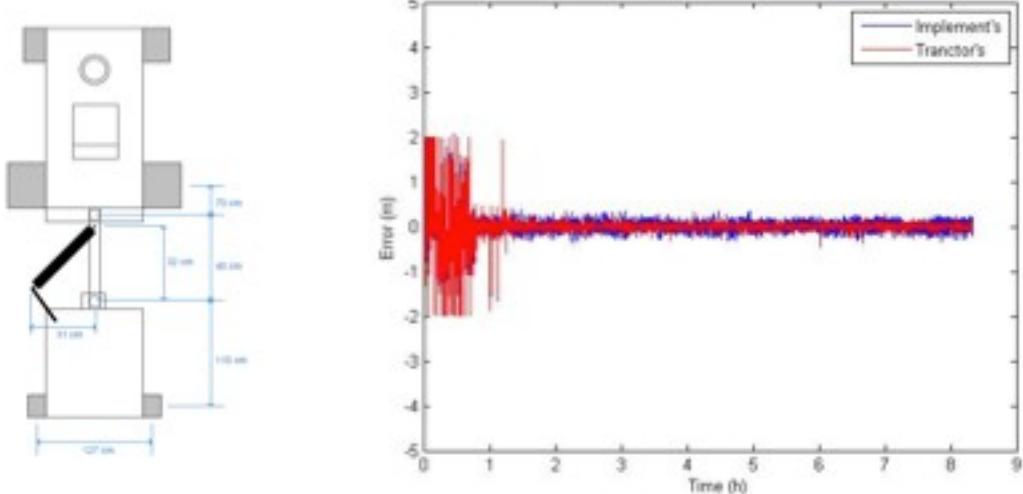
77

## Example: Autonomous tractor



78

## Results



79

## Conclusions

- Learning automata can be used as basic building blocks for more complex multi-state, multi-agent algorithms
- Game theoretic methods can still be used to analyze algorithms
- Methods can also be extended to continuous state-action spaces

80

## References

- [Kok et al, 2005] J. Kok, P. 't Hoen, B. Bakker, and N. Vlassis. Utile coordination: Learning interdependencies among cooperative agents. In Proceedings of the IEEE Symposium on Computational Intelligence and Games (CIG05), pages 29–36, 2005.
- [Melo & Veloso, 2009] F. Melo and M. Veloso. Learning of coordination: Exploiting sparse interactions in multiagent systems. In Proceedings of the 8th International Conference on Autonomous Agents and Multi-Agent Systems, pages 773–780, 2009.
- [De Hauwere et al, 2009] Y.-M. De Hauwere, P. Vranckx, A. Nowé, Multi-layer learning and knowledge transfer in MAS, In proceedings of the 7th European Workshop on Multi-Agent Systems, Ayia Napa, Cyprus, 2009.
- [De Hauwere et al, 2010] Y.-M. De Hauwere, P. Vranckx, A. Nowé, Learning Multi-Agent State Space Representations. In Proceedings of the 9th International Conference on Autonomous Agents and Multi-Agent Systems, pages 715–722, 2010.
- [De Hauwere et al, 2011] Y.-M. De Hauwere, P. Vranckx, A. Nowé, Solving sparse delayed coordination problems in multiagent reinforcement learning. Chapter in Adaptive Agents and Multi-agent Systems V, Springer-Verlag, 2011
- [Howell, 1997] Howell, M.N., Frost, G.P., Gordon, T.J. & Wu, Q.H. Continuous action reinforcement learning applied to vehicle suspension control. Mechatronics, xvii, 2, 3, 23, 24, 95, 1997
- [Rodriguez, 2011] A. Rodriguez, R. Grau, and A. Nowé. CONTINUOUS ACTION REINFORCEMENT LEARNING AUTOMATA. Performance and convergence. 3rd International Conference on Agents and Artificial Intelligence, pages 473–478. SciTePress, 2011.
- [Rodriguez, 2012] Rodriguez, A., Vranckx, P., Grau, R., & Nowé, A... Learning Approach to Coordinate Exploration with Limited Communication in Continuous Action Games. Proceedings ALA2012 Workshop (2012)

81

- [Rodriguez, 2013a] Rodriguez, A., Vranckx, P., Grau, R., & Nowé, A... A Reinforcement Learning Approach to Coordinate Exploration with Limited Communication in Continuous Action Games. Knowledge Engineering Review (in Press)
- [Rodriguez, 2013b] Rodriguez, A., Reinforcement Learning Approach to Coordinate Exploration with Limited Communication in Continuous Action Games. PhD thesis. Vrije Universiteit Brussel (2013)
- [Santharam, 1994] Santharam, G and Sastry, PS and Thathachar, MAL. Continuous action set learning automata for stochastic Optimization. In: Journal of the Franklin Institute, 331 (5). pp. 607-628, 1994
- [Verbeeck, 2004] Verbeeck, K. Coordinated Exploration in Multi-Agent Reinforcement Learning. PhD thesis, Vrije Universiteit Brussel, 2004
- [Vranckx et al, 2008] Vranckx, P., Verbeeck, K., & Nowé, A.. Decentralized Learning in Markov Games. (F.L. Lewis, Liu, D., & Lendaris, G.G.) IEEE Transactions on Systems, Man and Cybernetics (Part B: Cybernetics), 38, 976-81 (2008)
- [Vranckx et al, 2007] Vranckx, P., Verbeeck, K., & Nowé, A.. Optimal Convergence in Multi-Agent MDPs. Lecture Notes in Computer Science, Knowledge-Based Intelligent Information and Engineering Systems (KES 2007), 4694, (2007)
- [Vranckx et al, 2011] Vranckx, P., De Hauwere, Y.-M., & Nowé, A.. Transfer Learning for Multi-agent Coordination. In International Conference on Agents and Artificial Intelligence (ICAART). (2011)
- [Wheeler & Narendra, 1986] Wheeler Jr., R. & Narendra, K. Decentralized learning in finite markov chains. IEEE Transactions on Automatic Control, 31, 519–526. 37, 75, 1986
- [Witten, 1997] Witten, I.H. An adaptive optimal controller for discrete-time Markov environments. Information and Control, 34, 286–295. 37, 75, 1997

82